



# COMP 273

## The Control Unit (the sequencer)

Micro Architecture

Part 3

Prof. Joseph Vybihal



# Announcements

- Midterm exam next class
  - Crib sheets that I have included are:
    - Binary ASCII table
    - Classical CPU diagram
    - Pipeline CPU diagram
  - No crib sheets
  - Today's material is not on the exam




# Readings

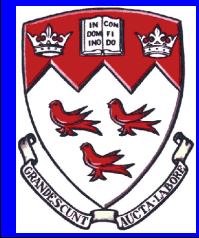
- Read
  - Edition 4 - Ch 4 sections 1 to 9
- Soul Of A New Machine
- Web Resources:
  - [http://www-static.cc.gatech.edu/classes/cs3760\\_99\\_spring/lectures/singlecycle.html](http://www-static.cc.gatech.edu/classes/cs3760_99_spring/lectures/singlecycle.html)



# Outline

- The Sequencer (control unit)
- Hazards, faults, stalls and dumps
- Calculating CPU performance





# Part 1

## The Control Unit

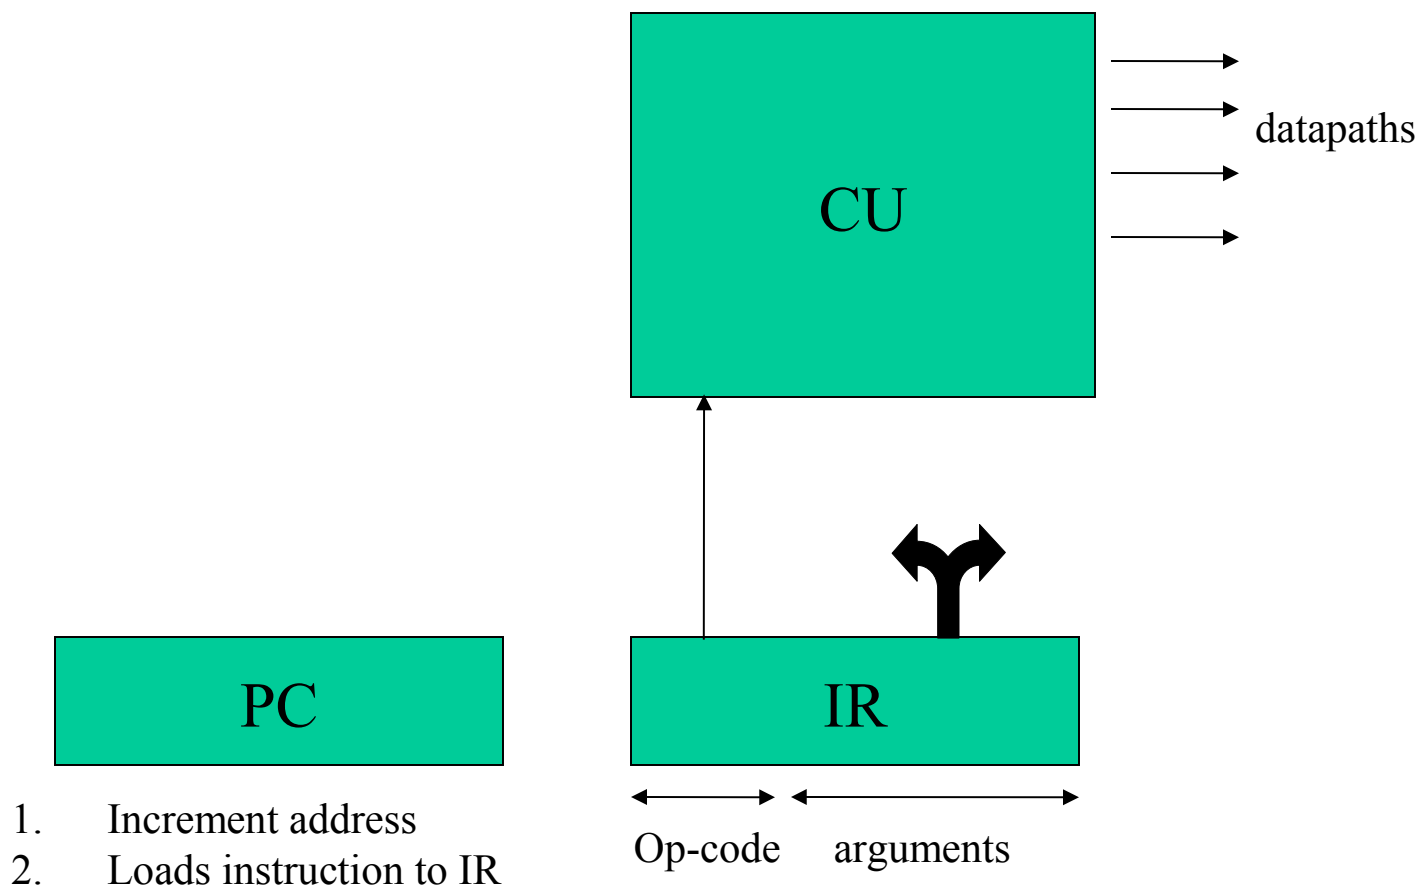


# Definitions

- Datapaths
  - The “wires” of a CPU that need to be engaged at a particular order & tick of the clock to implement an instruction.
  - This includes what the path is connected to: registers, gates, ALU, ...
- Control
  - That portion of the CPU, often called the control unit or sequencer, that is responsible for the timing and triggering of a datapath
- Instruction Format
  - The organization of bits in the IR (definition of an instruction)
- Micro-programming (the datapaths that implement the instruction)
  - Flat: One instruction executes at a time
  - Pipeline: Assembly execution of more than 1 instruction
  - Cores: Parallel execution




# Flat Control Unit





# Flat Control Unit

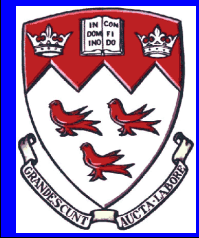


Counter logic

For instructions that require  $n$  ticks

Tick / tick  
Megam

© 2015







Classical:

Add r3, r2, r1

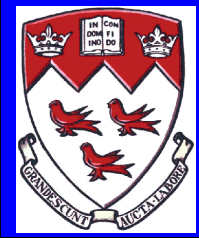
Tick: L  $\leftarrow$  r2

Tick: R  $\leftarrow$  r1

Tick: A  $\leftarrow$  add

Tick: r3  $\leftarrow$  A

4 steps






# Op-code and Count

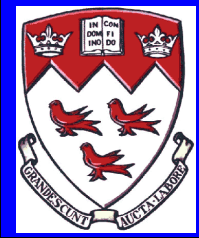
OP-CODE	COUNT
00100	00100
	00101
	00110
01000	01000
	01001
	01010
	01011

01100



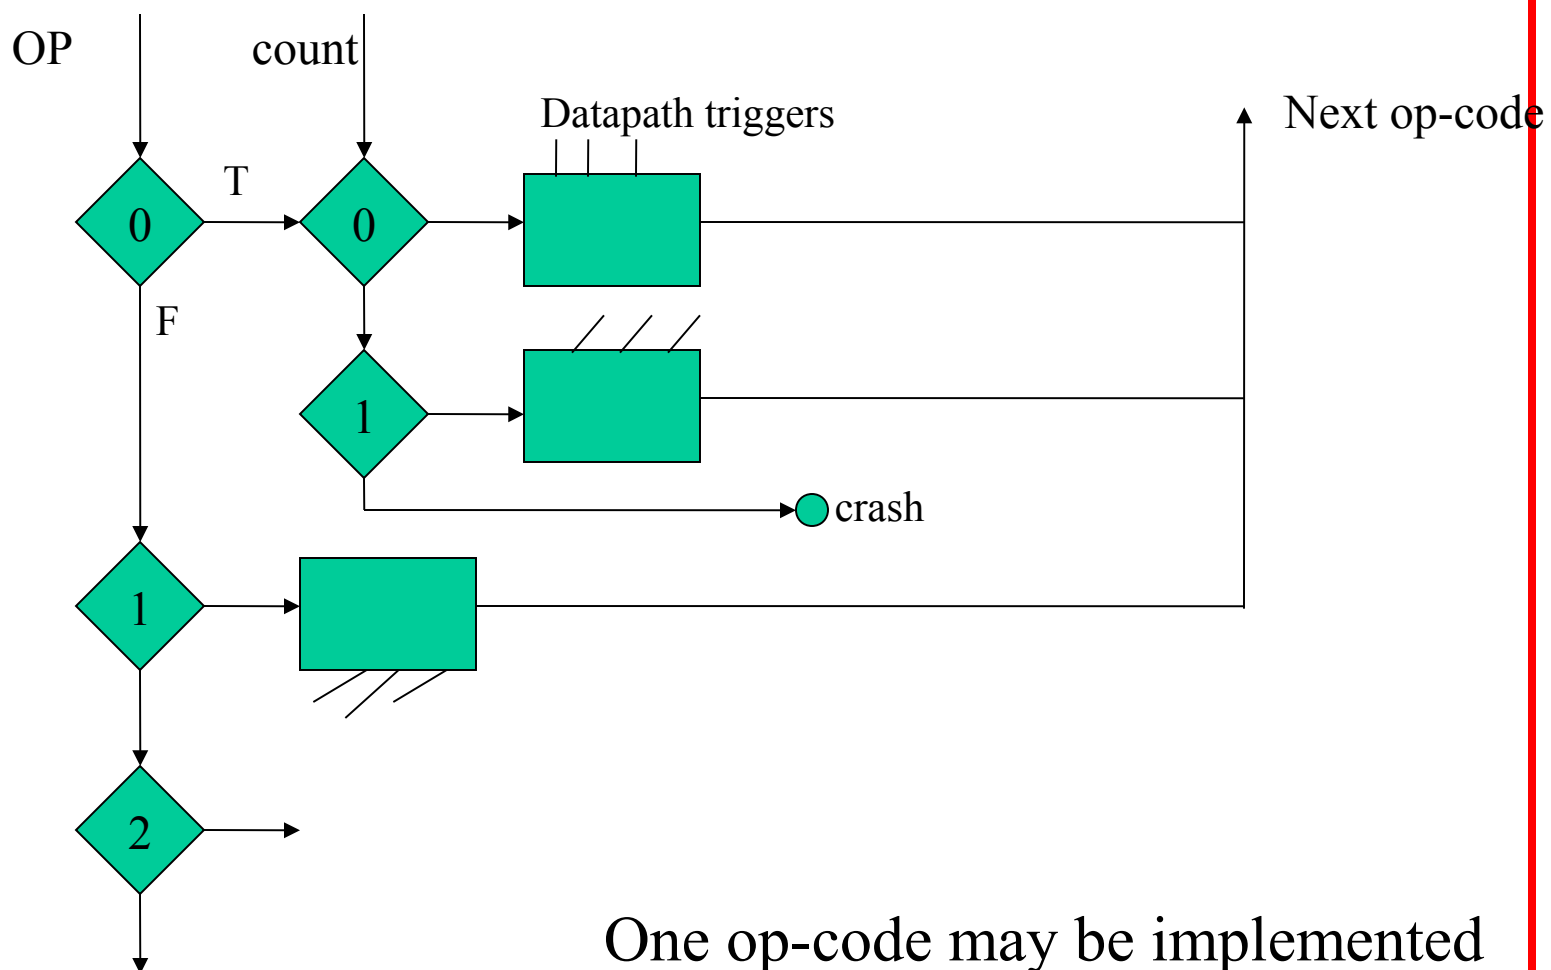
Actual op

The count (built into op-code or sign extended in register)

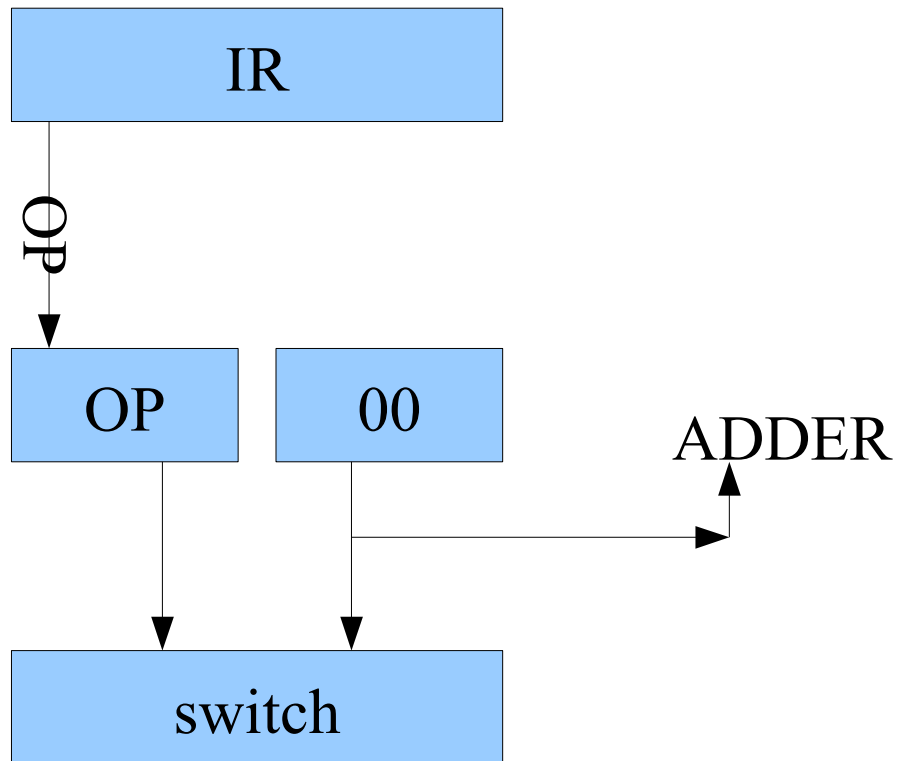
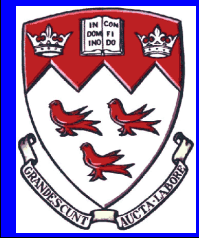




# Flat Sequence



One op-code may be implemented via multiple sub-steps (datapaths)



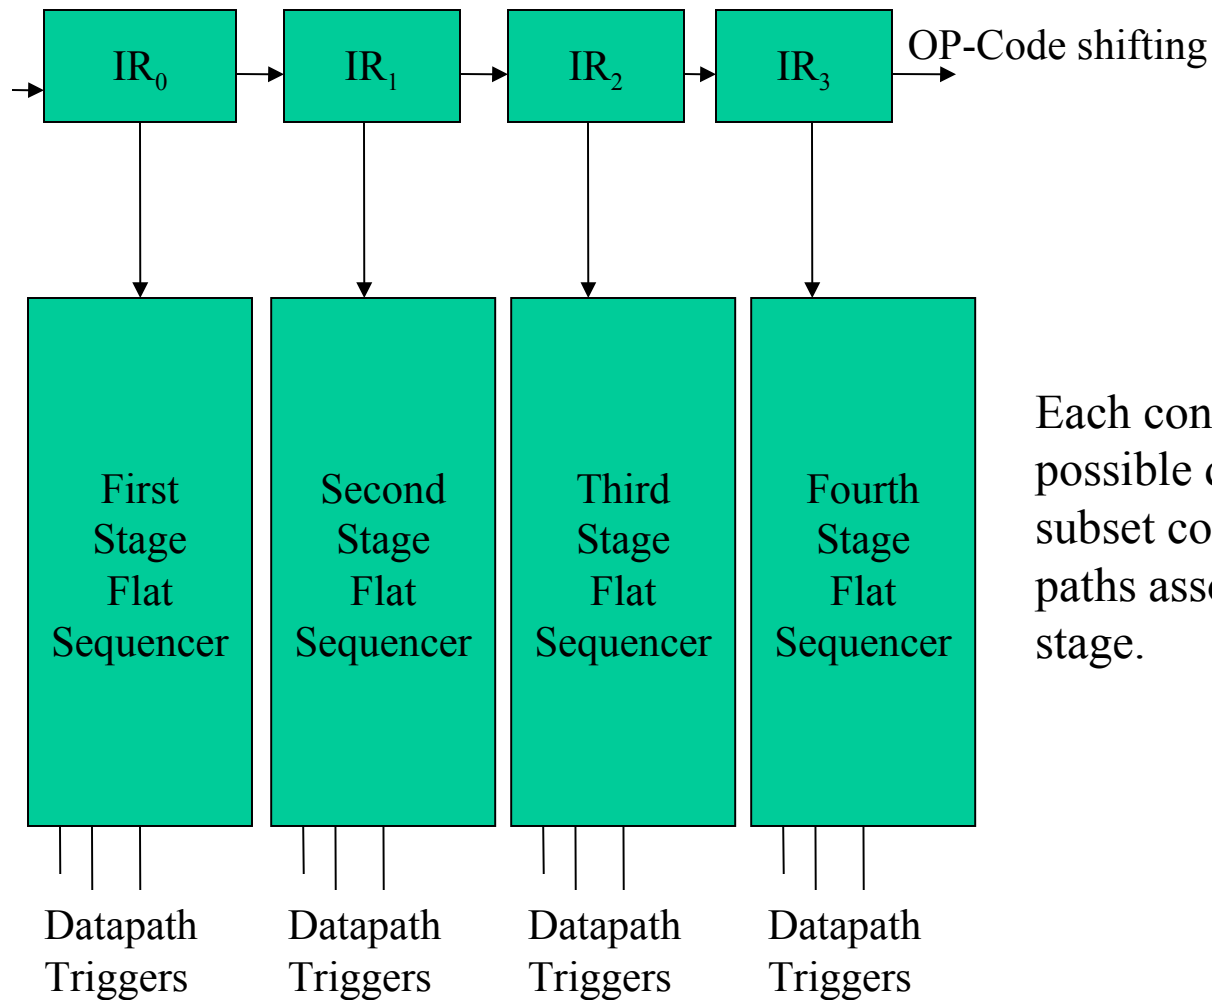
Case 0: data 0

Case 1: data 1



4 IR buffers


# Pipeline Sequence



Each contains a subset of all possible datapaths. This subset consists of only those paths associated with the stage.

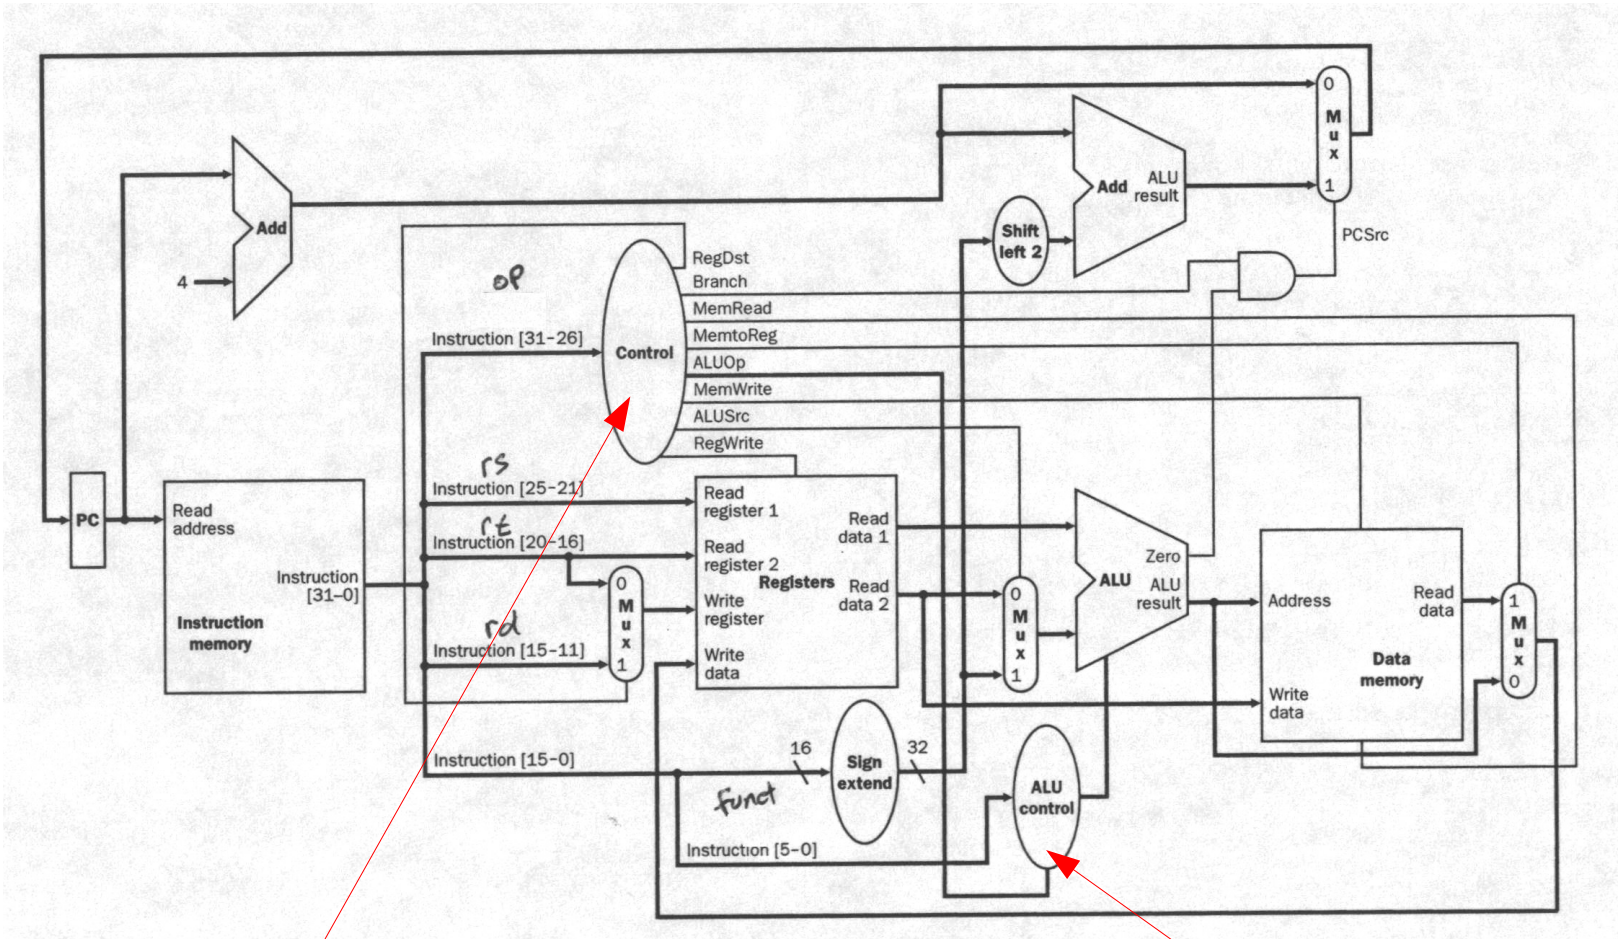


# Pipeline Sequencer (basic)






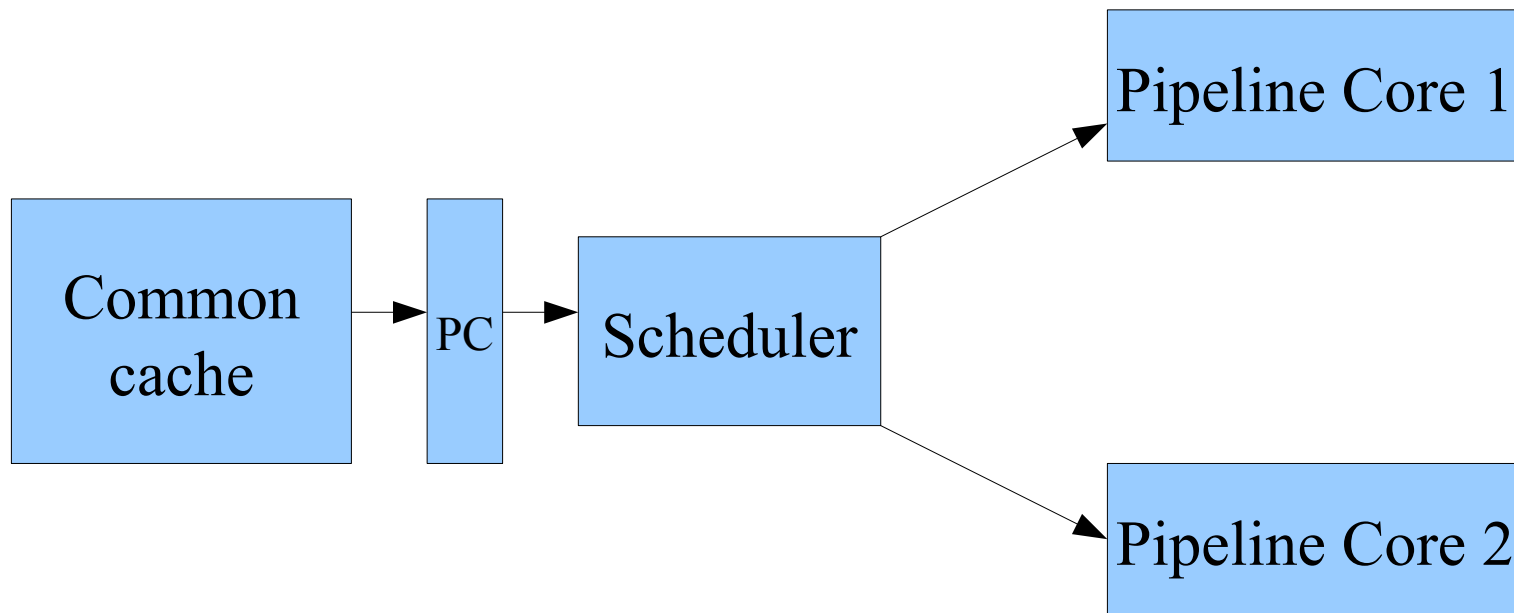
# MIPS Sequencer



Two sequencer units



# Cores Sequencer (simplified)



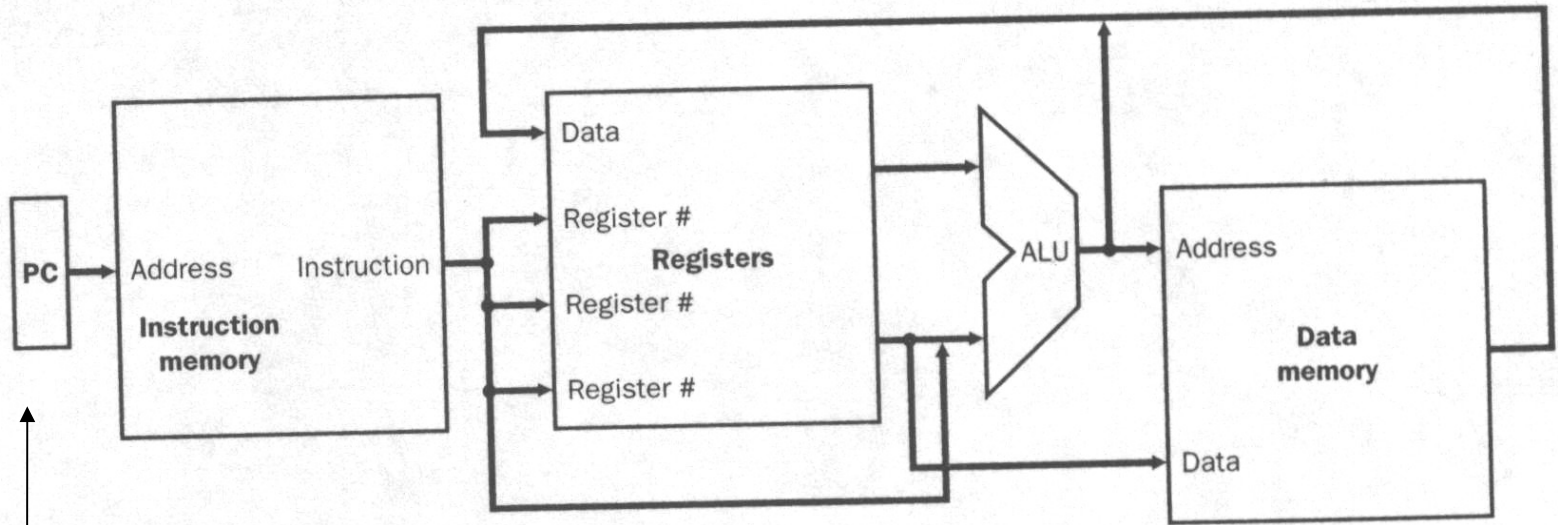




# Buffers




# Buffers and Multiple Pathways



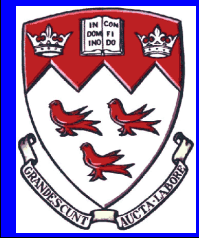
One PC

Multiple IR registers (buffers) for each instruction in pipeline, each passing through one stage of the control unit (sequencer)





IR (+ buffers)

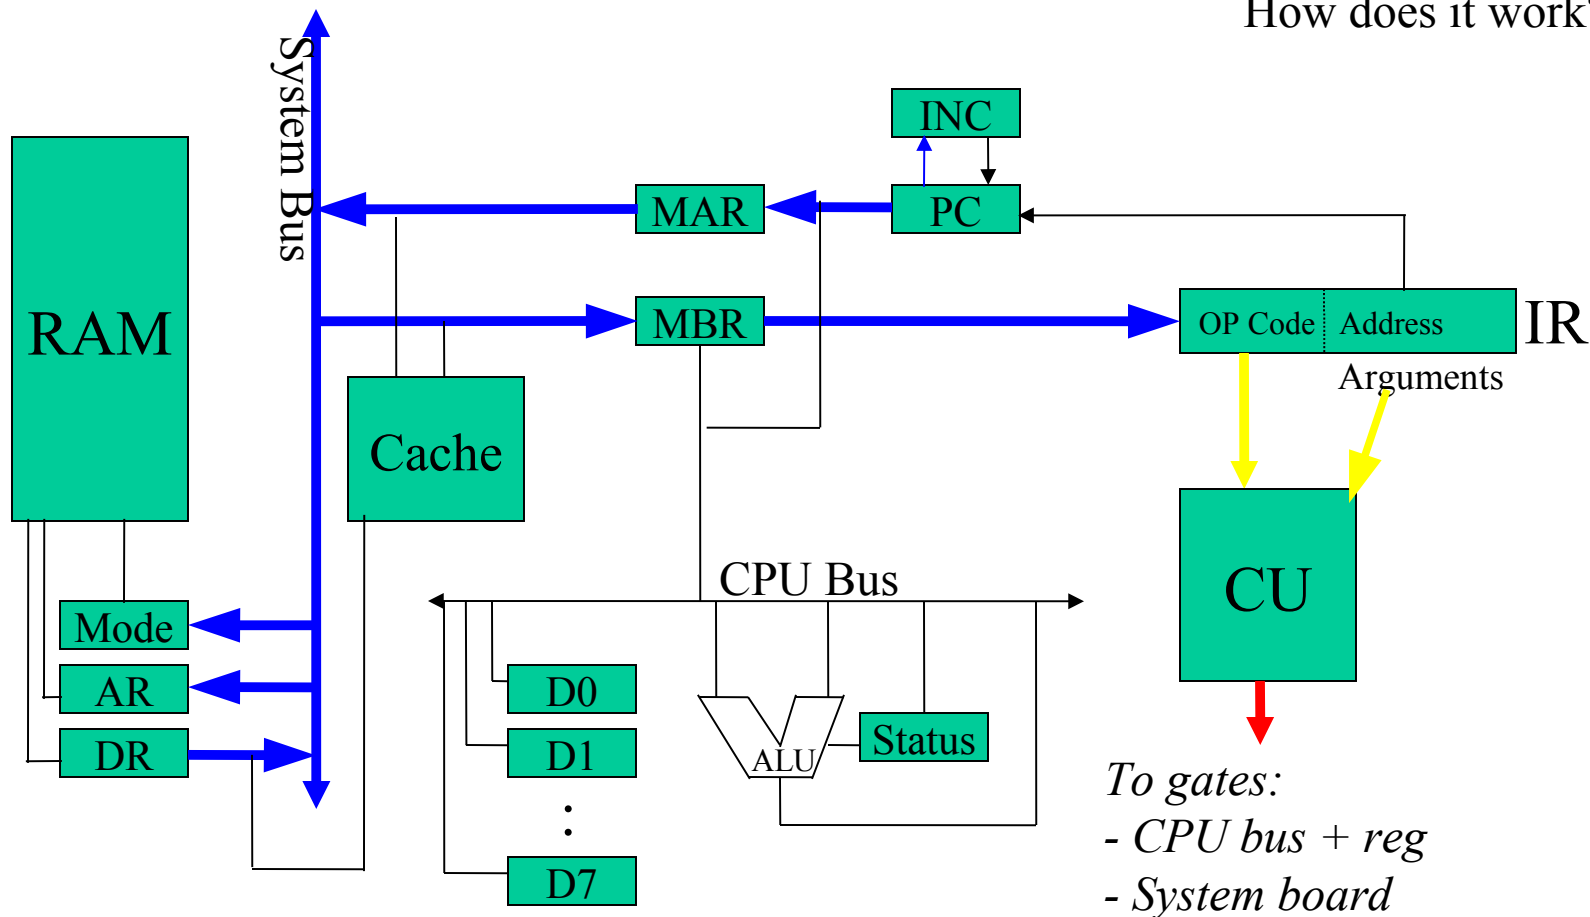


# CPU Execution Overview/Review



# Classical CPU Design

How does it work?



To gates:  
- CPU bus + reg  
- System board

**1. Fetch**  
**(5 cs)**

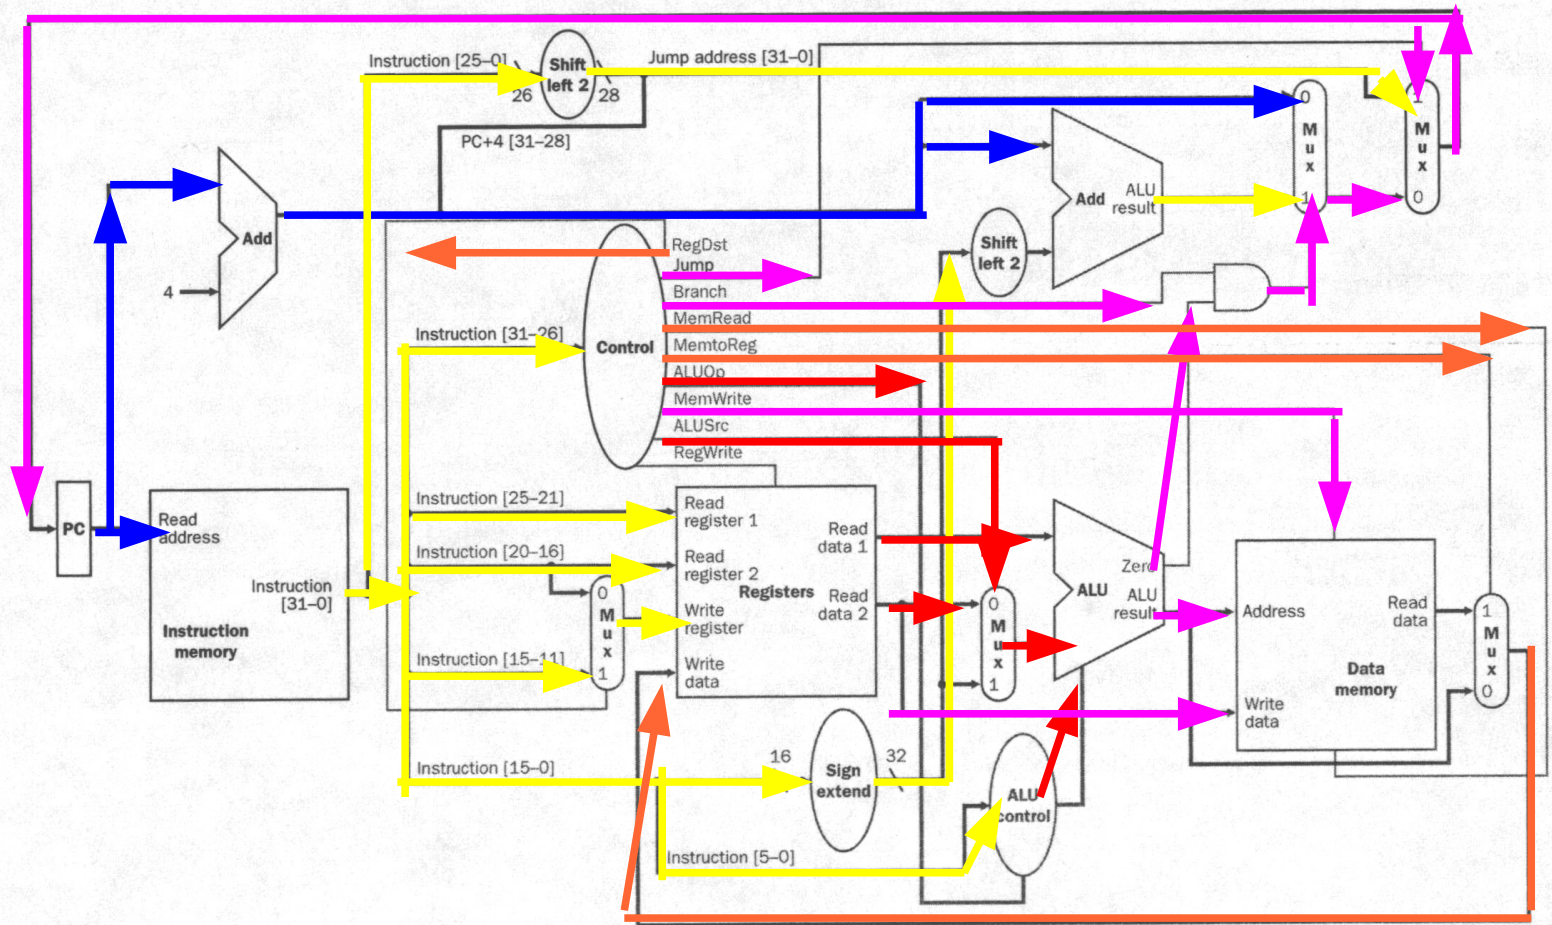
**2. Decode**  
**(1 cs)**

**3. Execute**  
**(1 cs per step = 3 – 5 steps)**

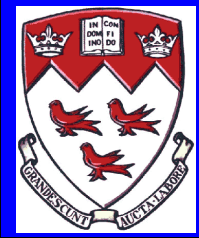


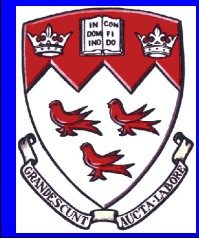


# Pipeline CPU Design



1. Fetch (1 cs)
2. Decode (1 cs)
3. Execute (1 cs per step)
4. Store





# Hazards & Faults



# Terminology

- A “hazard” may lead to a “fault”
  - Hazard: a danger to keep watch for
  - Fault: an error that has occurred
- A “fault” leads to a CPU execution “stall”
  - Stall: Normal CPU execution cycle lengthens
- A “stall” can lead to a “dump” or a “no-op”
  - Dump: All instruction after the fault are dumped from the pipeline and re-loaded after fault is handled (major slow down)
  - No-Op: Additional instruction(s), that do no action, inserted between the fault and the next instruction to allow the CPU to execute the problematic instruction without side-effects (minor slow down)



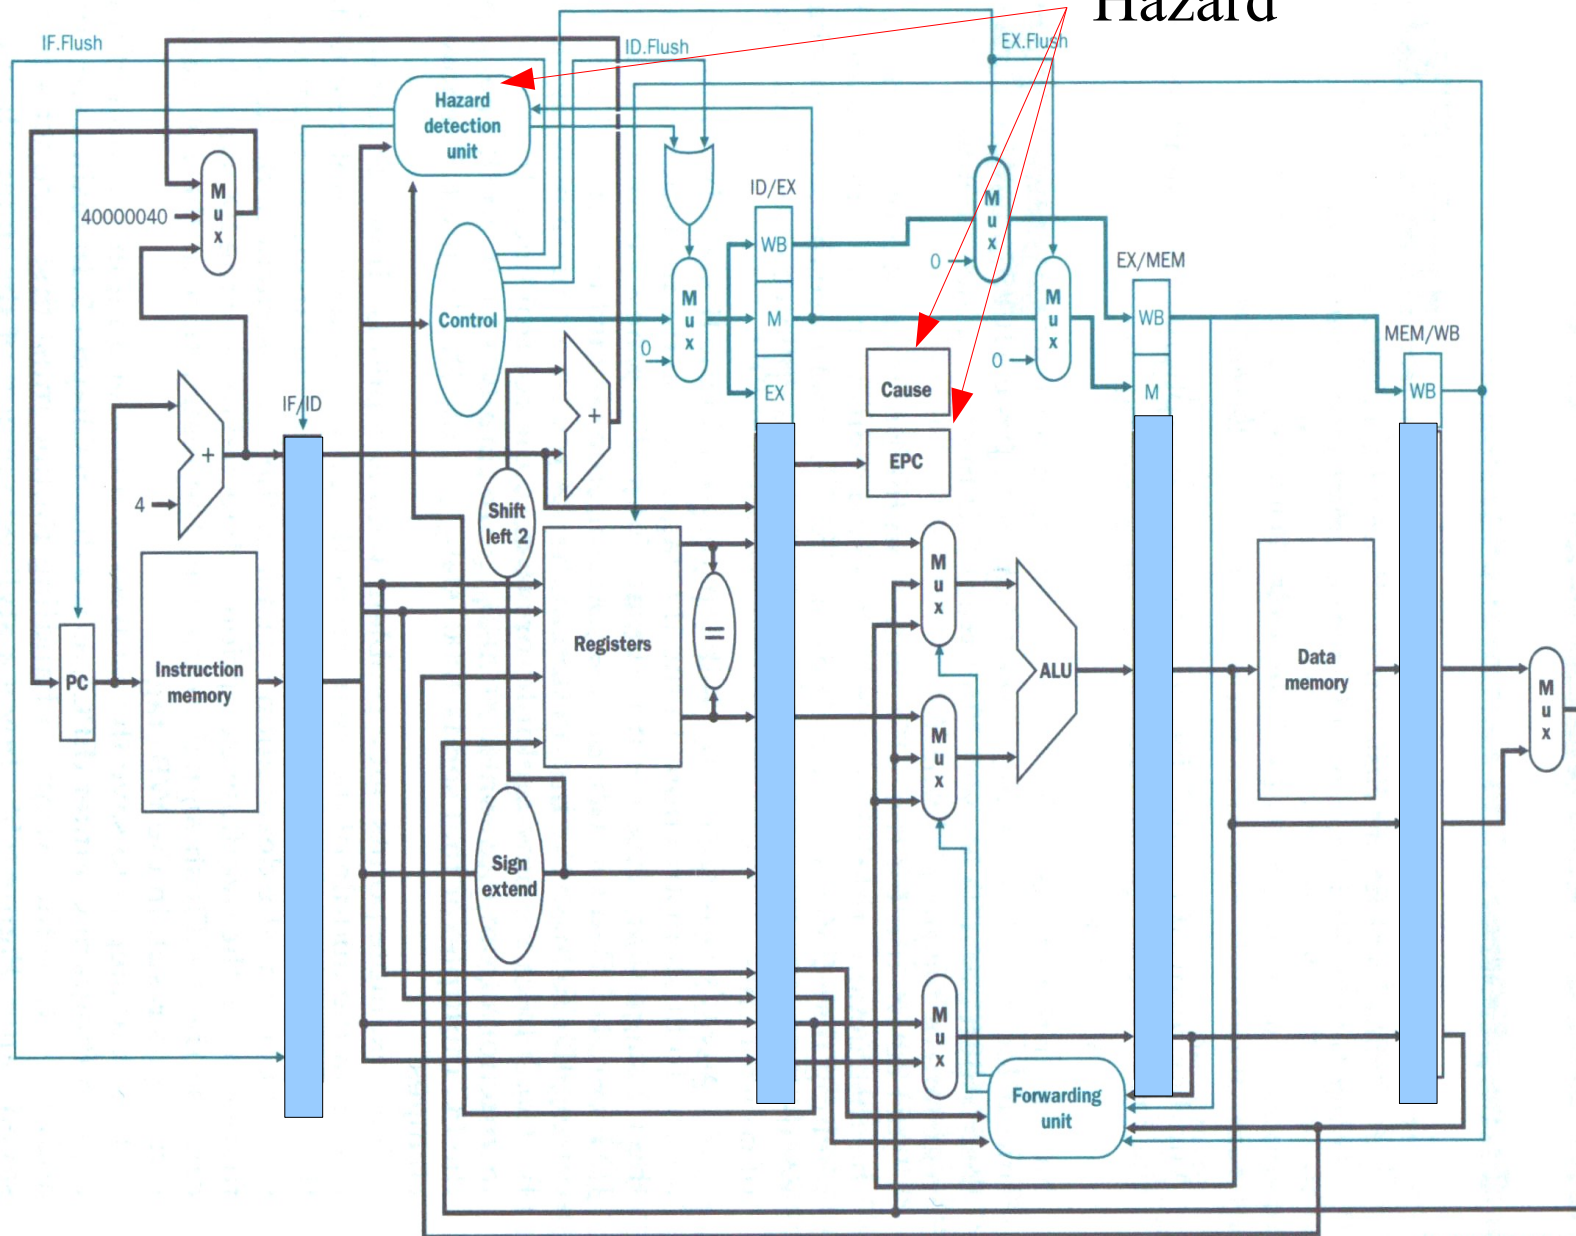


# Types of Hazards

- Hazards
    - Structural Hazards
      - CPU cannot support the combination of instructions in the pipeline (eg. Single instruction store/load crash)
      - Illegal instruction or illegal result (like divide by zero)
    - Control Hazards
      - Branch request causing the semi executed pipeline instructions to be unnecessary
    - Data Hazards
      - An instruction depends on the results of a previous instruction in the pipeline
        - Add \$s0, \$t0, \$t1
        - Sub \$t2, \$s0, \$t3
- Identifiable by bit patterns
- NOP or wiring for Forwarding or Backup buffer

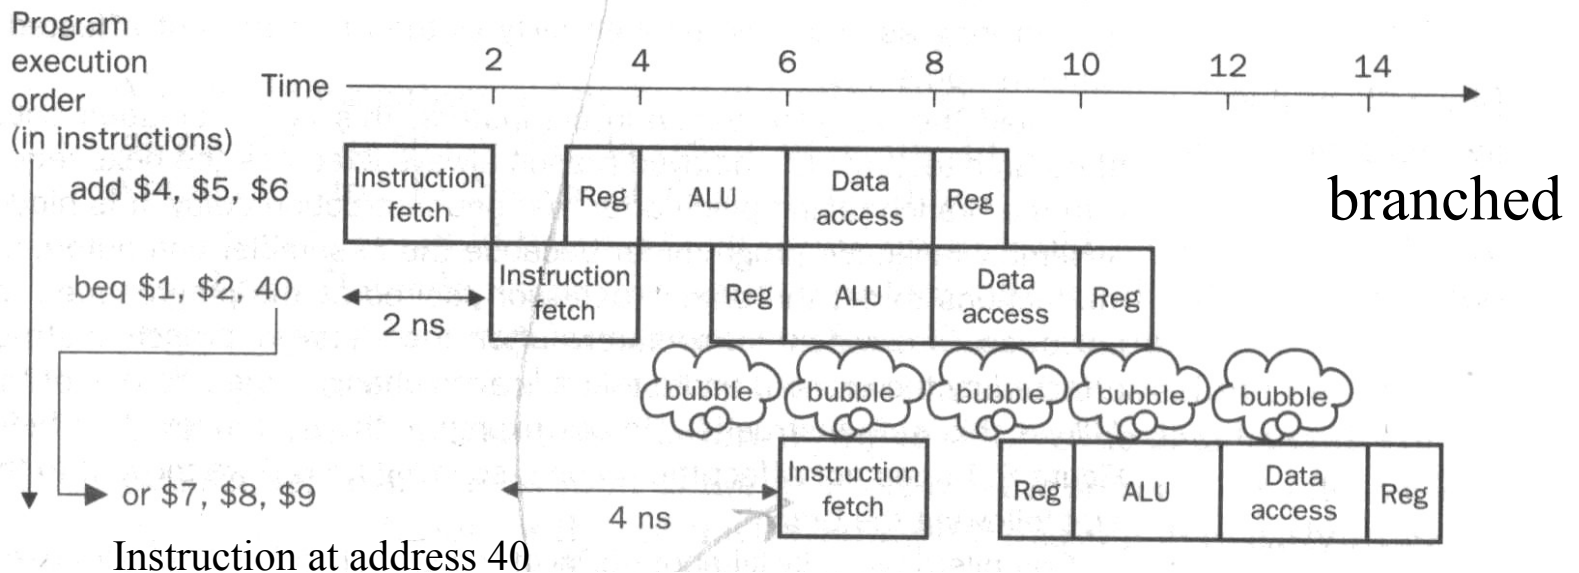
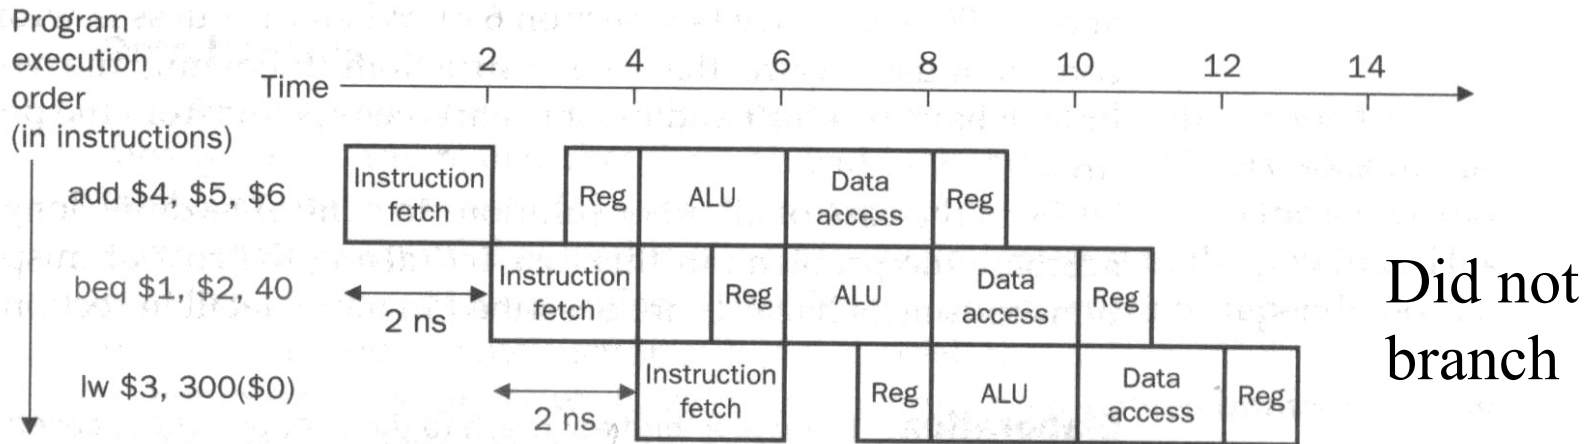


# Hazard






# Effects...





# Calculating Loss

- At which stage did the fault occur?
- If no-op, how many were inserted?
- If dump, how many instructions need to be reloaded?




Dump  
Points



# Calculating Loss

- How many instructions execute at full speed (no faults)?
- What do we lose when we do a stage 2 (IR1), stage 3 (IR2) dump or stage 4 (IR3) branch?





# Calculating Loss

- What performance loss do we experience?
  - $\text{Expected}(\text{Loss}) = P(\text{Loss}) * \text{Cost}$
- How often can we expect a stall?
  - 17% on average
    - Actually depends on the program,





# Solutions?


Inserting no-op(s) or reordering

- Prediction
  - Assume branch will always fail (not used)
    - Happens too often...
  - Assume branch success based on type
    - Function calls (yes)
    - Backward jumps (probably loops, yes)
    - If statements (no, fail)
  - Reorder instructions (compiler solution)
    - We need a delay to figure out if branch will happen
    - Instead of using a NOP, reorder branch & preceding instruction



# Reorder Example

- Original code:
  - Add \$4, \$5, \$6
  - Beq \$1, \$2, 40
  - Lw \$3, 300(\$0)



- Reorder:
  - Beq \$1, \$2, 40
  - Add \$4, \$5, \$6
  - Lw \$3, 300(\$0)

Finishes at the same time

← This is not loaded





# Part 2

## CPU Performance Issues



# Cycles vs Clock Ticks

- Clock Ticks
  - A count of the number of actual clock ticks required to perform an activity in the CPU
- Cycles (confused definitions)
  - A count of the number of micro instruction required to perform an activity in the CPU
- Note: often equations with cycles can be used with clock ticks



# The Role of Performance

Keep in mind that the only complete and reliable measure of CPU performance is **clock ticks**. *CPU Execution Time* is the product of three basic measurements:

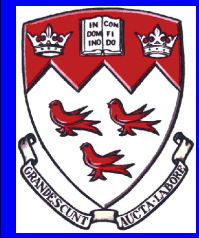
- Instruction count (total number of instructions in program)
- Clock cycles per instruction
- Clock cycle time

CPU execution time in seconds =  
X instructions in program \*  
Y ticks per instruction \*  
Z seconds per tick



# Example 1

- A 2 Ghz computer ticks 2 billion times/sec.
- If we have CPU that executes a single instruction in a total of 20 ticks
- How long will it take to run a program that has 1000 lines of code (linear)?
  - $T = 1000 * 20 / 2\text{billion} = 0.00001 \text{ sec}$





# Example 2

Assume that the operation times for the major functional units in a CPU are:

- Memory access: 2 ns
- ALU and adders: 2 ns
- Register file (i/o): 1 ns

Assuming that the multiplexers, control unit, etc. have no delay, which of the following implementations would be faster and by how much?

1. An implementation in which every instruction operates in 1 clock cycle of a fixed length.
2. An implementation where every clock cycle is of variable length equivalent to the instruction step to execute

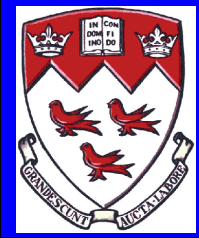


Assuming the program instruction mix is:

- 24% loads
- 12% stores
- 44% ALU
- 18% branches
- 2% jumps

Recall:

CPU Time =  $X$  instructions \*  $Y$  ticks/instr \*  $Z$  seconds/ticks





Instruction class	Functional units used by the instruction class				
ALU type	Instruction fetch	Register access	ALU	Register access	
Load word	Instruction fetch	Register access	ALU	Memory access	Register access
Store word	Instruction fetch	Register access	ALU	Memory access	
Branch	Instruction fetch	Register access	ALU		
Jump	Instruction fetch				

2

1

2

1 or 2

1

Memory access: 2 ns

ALU and adders: 2 ns

Register file (i/o): 1 ns

Instruction class	Instruction memory	Register read	ALU operation	Data memory	Register write	Total
ALU type	2	1	2	0	1	6 ns
Load word	2	1	2	2	1	8 ns
Store word	2	1	2	2		7 ns
Branch	2	1	2			5 ns
Jump	2					2 ns

Variable vs. fixed cycle lengths?



Instruction class	Instruction memory	Register read	ALU operation	Data memory	Register write	Total
ALU type	2	1	2	0	1	6 ns
Load word	2	1	2	2	1	8 ns
Store word	2	1	2	2		7 ns
Branch	2	1	2			5 ns
Jump	2					2 ns

CPU Time fixed = longest instruction = load = 8 ns

CPU time variable =  $8 * 24\% + 7 * 12\% + 6 * 44\% + 5 * 18\% + 2 * 2\% = 6.3$  ns

Performance improvement =  $8 / 6.3 = 1.27$

Variable clock cycles are not always used:

- Harder to build
- Delay in logic gates to implement  $> 1.27$