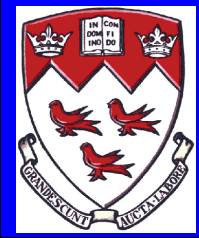




COMP 273

Cache Memory

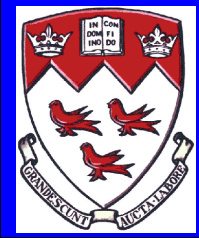
Prof. Joseph Vybihal





Announcements

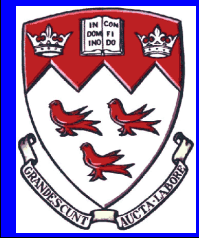
- Course evaluation





At Home

- Using the supplied code with the simulator, manipulate the cache.
- Textbook:
 - See MIPS Run; By Sweetman; Morgan Kaufmann Publishers, ISBN 1-55860-410-3 Chapter 4



Part 1

Cache Basics



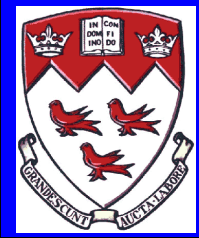
Why use a cache?

- To speed up the operation of the computer
- Primary problem: it is smaller than the program, wouldn't you spend more time in the RAM?



Principle of Locality

- Temporal Locality
 - Item will be referenced again soon
 - Example: Library and functions
- Spatial Locality
 - Adjacent items will probably be executed next
 - Example: Loops and functions





The Storage Conundrum


- We need fast and voluminous storage to run our applications adequately
- Two issues at hand:
 - Speed of access
 - Amount of storage
- Speed and storage physically work against each other, when compared to cost in dollars



Storage Comparison

Storage	Technology	Speed	Cost
CPU Registers	Flip-flops	1 – 5 ns	\$250 - \$300
Cache	SRAM – transistors + power	5 – 25 ns	\$100 - \$250
RAM	DRAM – capacitors + refresh	30 – 120 ns	\$5 - \$10
Disk	Magnetic charge – mechanical	10 Million ns – 20 million ns	\$0.1 - \$0.2

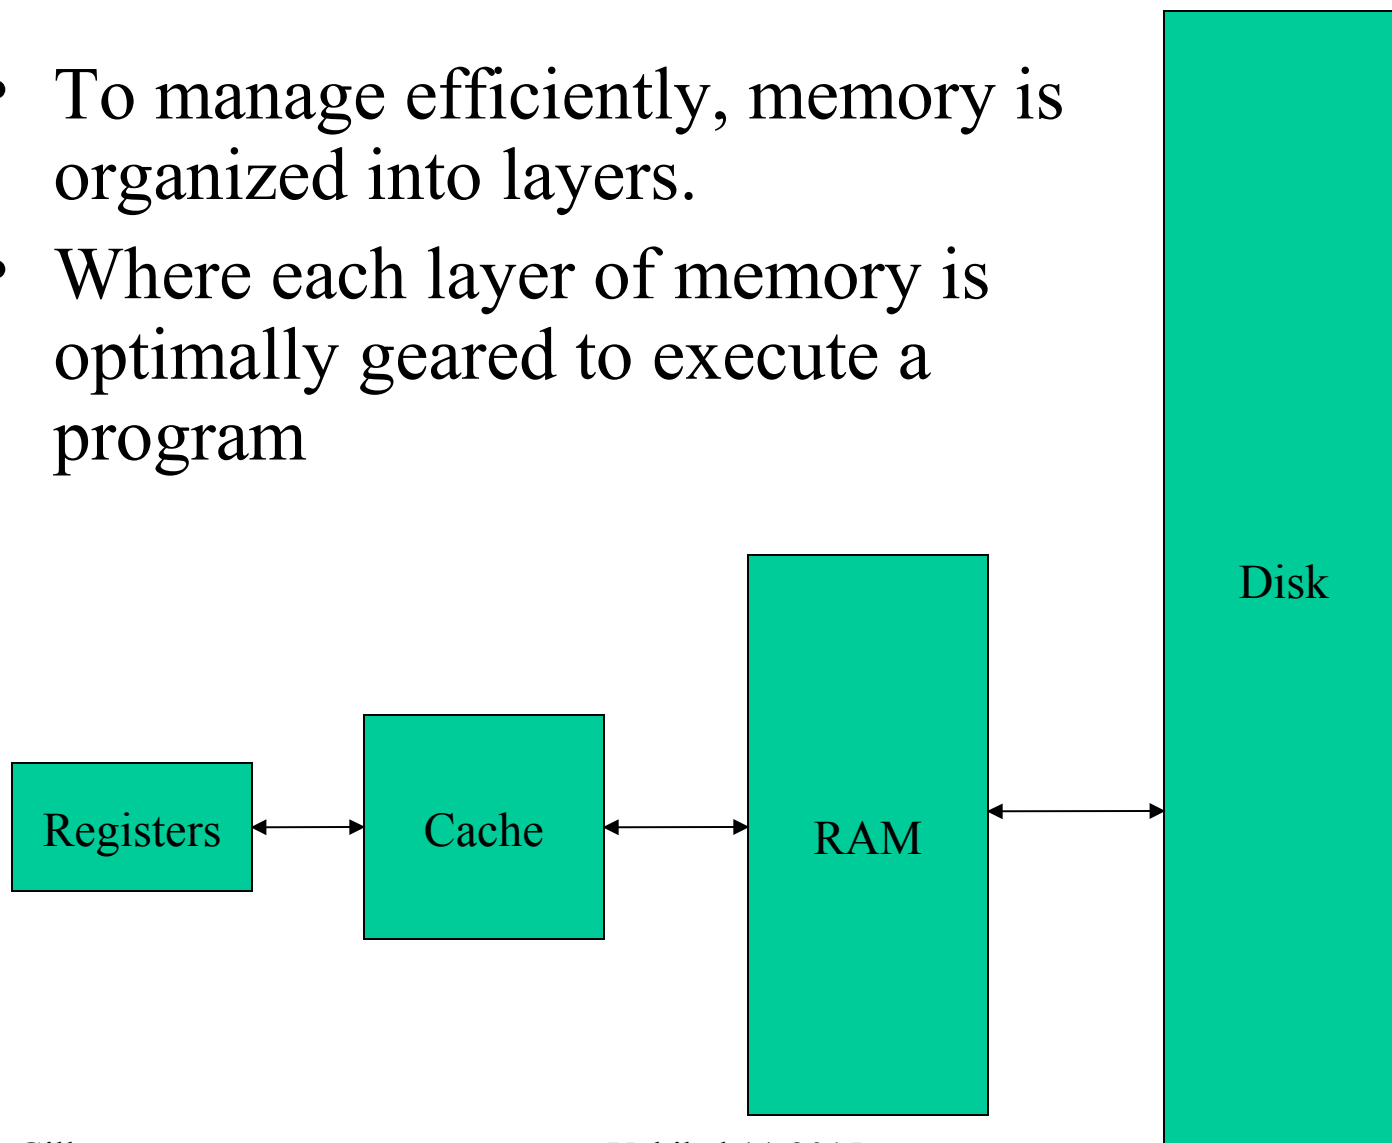
~ Per 100 Meg





Memory Hierarchy

- To manage efficiently, memory is organized into layers.
- Where each layer of memory is optimally geared to execute a program





Cache Levels

TABLE 4.1 Cache evolution in MIPS CPUs

CPU (MHz)	Primary				Secondary			Tertiary		
	Size I-cache	D-cache	direct/ n-way	on- chip?	Size	direct/ n-way	on- chip?	Size	direct/ n-way	on- chip?
R3000-33	32K	32K	Direct	Off	1M	Direct	Off	8M	Direct	Off
R3052-33	8K	2K	Direct	On						
R4000-100	8K	8K	Direct	On						
R4600-100	16K	16K	Two-way	On						
R10000-250	32K	32K	Two-way	On						
R5000-200	32K	32K	Two-way	On						
RM7000-xxx	16K	16K	Four-way	On						





How to optimally use cache?

Problems ...

- Issues:
 - What to load?
 - Can we make an educated guess?
 - How to address?
 - Should we match with RAM's address?
 - Hit to Miss ratio (cache miss rate)
 - Hit implies we find the instruction in the cache
 - Miss implies we need to go to the RAM
 - Cache miss/refill penalty



Operations in a Miss

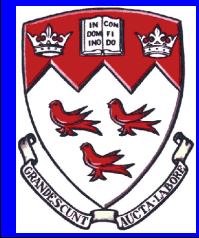
1. Access RAM (PC – 4, or jump address)
2. Wait for RAM to complete read
3. Put the data into cache & update table
4. Restart the instruction (from cache)

Miss penalty = Cycles to upload data to cache

Cost of miss = Miss frequency * Miss penalty

Program speed = $n + m * \text{penalty}$, where n & m are no. of instructions

(compare with no cache?)



Part 2

Cache Mechanics




Addressing

Block address MODULO Number of cache blocks in the cache

$$\log_2 8 = 3 \text{ bits}$$

- Lower bit mapping
- Not 1:1





Basic Structure & Operation


E.g.

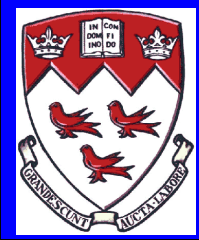
Decimal address of reference	Binary address of reference	Hit or miss in cache	Assigned cache block (where found or placed)
22	10110_{two}	miss (7.6b)	$(10110_{\text{two}} \bmod 8) = 110_{\text{two}}$
26	11010_{two}	miss (7.6c)	$(11010_{\text{two}} \bmod 8) = 010_{\text{two}}$
22	10110_{two}	hit	$(10110_{\text{two}} \bmod 8) = 110_{\text{two}}$
26	11010_{two}	hit	$(11010_{\text{two}} \bmod 8) = 010_{\text{two}}$
16	10000_{two}	miss (7.6d)	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$
3	00011_{two}	miss (7.6e)	$(00011_{\text{two}} \bmod 8) = 011_{\text{two}}$
16	10000_{two}	hit	$(10000_{\text{two}} \bmod 8) = 000_{\text{two}}$
18	10010_{two}	miss (7.6f)	$(10010_{\text{two}} \bmod 8) = 010_{\text{two}}$

MAP IN unique DATA

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory(10110_{two})
111	N		

Cache Memory





Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

a. The initial state of the cache after power-on

Index	V	Tag	Data
000	N		
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

c. After handling a miss of address (11010_{two})

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	Y	00 _{two}	Memory (00011 _{two})
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

e. After handling a miss of address (00011_{two})

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory(10110 _{two})
111	N		

b. After handling a miss of address (10110_{two})

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		



d. After handling a miss of address (10000_{two})

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	10 _{two}	Memory (10010 _{two})
011	Y	00 _{two}	Memory (00011 _{two})
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

f. After handling a miss of address (10010_{two})




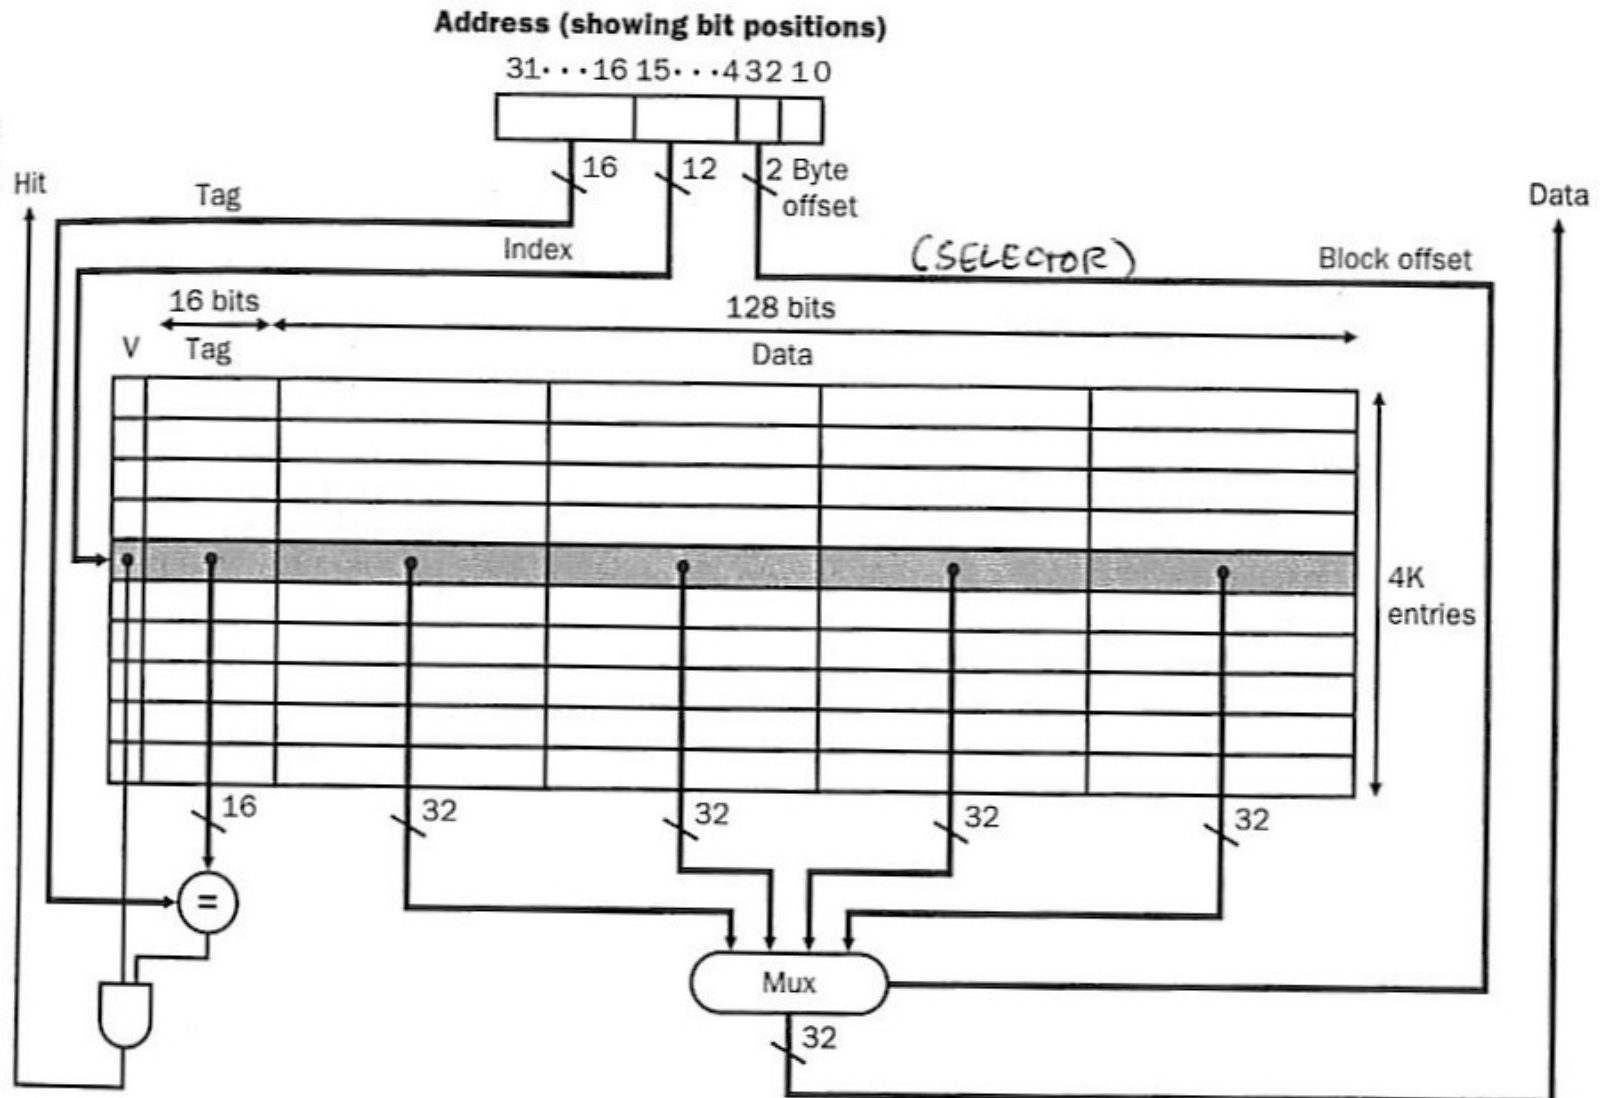
Basic Access Architecture





Locality Access

(Read/Miss = 4 word read, a block read)






Performance

Program	Block size in words	Instruction miss rate	Data miss rate	Effective combined miss rate
gcc	1	6.1%	2.1%	5.4%
	4	2.0%	1.7%	1.9%
spice	1	1.2%	1.3%	1.2%
	4	0.3%	0.6%	0.4%

Just temporal


Spatial locality advantage






COMP 273

Introduction to Computer Systems



Experiments on cache sizes


← MAGIC #





Part 3

Technology to support caches





Bus Designs

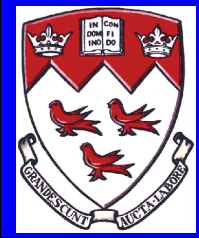
- One-Word-Wide memory organization
 - $\text{RAM} \rightarrow \text{Bus} \rightarrow \text{Cache} \rightarrow \text{CPU}$
- Wide memory organization
 - $\text{RAM} \rightarrow \text{multi-word bus} \rightarrow \text{Cache} \rightarrow \text{Multiplexor} \rightarrow \text{CPU}$
- Interleaved memory organization
 - $\text{RAM banks} \rightarrow \text{Bus} \rightarrow \text{Cache} \rightarrow \text{CPU}$

Assume:

- 1 cs to send an address to RAM
- 15 cs for each DRAM access find
- 1 cs to send a word of data


What improvement
for 4 word lw?

(one instruction but 4 word load)



Part 4

Performance Calculations



Amdahl's Law



Amdahl's Law

- The speedup in performance is proportional to the new component and the actual fraction of work it carries out.

$$s = 1 / [(1 - f) + f/k]$$

Where:

- S is the speedup
- F is the fraction of work performed by the component
- K is the advertised speedup of the new component



Example

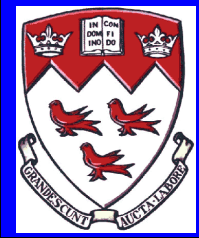
- Assume your daytime processes spend 70% of their time running in the CPU and 30% waiting for service from a disk.
- You find a computer that functions 50% faster and costs \$ 10,000.
- You find a new disk drive for \$7,000 with a speed increase of 2.5
- What should you do?

Processor: $f = .7$, $k = 1.5$

$S = 1 / [(1 - 0.7) + 0.7/1.5] = 1.3 \rightarrow \underline{30\%}$ for \$10,000 = $10000 / 30 = \$333$ per

Disk: $f = .3$, $k = 2.5$

$S = 1 / [(1 - 0.3) + 0.3/2.5] = 1.22 \rightarrow 22\%$ for \$7,000 = $7000 / 22 = \underline{318}$ per



Transfer Rate Calculations



Polling Overhead

- Assume polling overhead takes 400cs on a CPU that runs at 500Mhz
- How much CPU time is used:
 - Poll a mouse 30 times per second
 - Floppy disk data transfer rate at 16-bits per tick and needs to move data at a rate of 50KB/sec.
 - Hard disk transfers at 4-word (32 bits per) chunks at 4MBytes/sec.



Answer

- Mouse
 - Polling = $30 * 400 = 12,000cs$
 - Processor = $12000 / 500,000,000cs = 0.002\%$
- Floppy Disk
 - Polling = $50 / 2 = 25K$ times
 - $25K * 400 = 10,000,000cs$
 - Processor = $10 / 500 = 2\%$
- Hard Disk
 - Polling = $4MB / 16Bytes = 250K$ times
 - $250K * 400 = 100,000,000$
 - Processor = $100 / 500 = 20\%$



Interrupt Driven

- Assume 500MHz computer 4MB/s HDD 4-words
- Assume 500cs to handle interrupt
- Hard disk used 5% of the time
- Processor usage?

Data transfer = $4\text{MB} / 16\text{B} = 250\text{K}$ times

(interrupt at each of these moments, like in polling)

Polling = $250\text{K} * 400 = 100,000,000\text{cs}$

Interrupt = $250\text{K} * 500 = 125,000,000\text{cs} \quad !!$

Polling / Processor = $100 / 500 = 20\%$ (always)

Interrupt / Processor = $125 / 500 = 25\%$ (at transfer)

Interrupt used only 5% of the time = $25\% * 5\% = 1.25\%$ (actual)

Interrupt advantage over polling occurs during the time when no data transfer occurs (absence of overhead).



DMA

- Assume 500MHz, 4MB/sec DMA 1-byte
- DMA overhead (initialize) 1000cs
- Interrupt overhead (process at end) 500cs
- DMA transfer rate of 8 KB per cycle
- Processor usage?

DMA transfer = 8KB / 4MB/sec = 0.002 seconds (w/o CPU)

100% of the time = (1000 + 500) = 1500 cs

Processor = 1500 / 500000000 = 0.000003 => 0.0003%

Only for overhead ... no transfer work