



# Distributed Transactions


(a) Centralized




(c) Middleware approach



(b) Client-Aware Distribution




(d) Homogenous Approach



COMP-512: Distributed Systems

# Distributed Transaction Processing

- ❑ Collection of sites  $S^i, i = 1, \dots, N$
- ❑ Each process is responsible for a (finite) set  $D^i$  of data items;
- ❑ We first assume no replication, that is  $D^i \cap D^j = \emptyset$  for  $i, j = 1, \dots, N, i \neq j$
- ❑ Homogenous: each server knows the location of all data items
- ❑ Middleware: middleware knows the location of all data items
- ❑ Transaction  $T_i$  is submitted to one site  $S^l$ , called the coordinator of  $T_i$  (or to the global transaction manager / middleware)
  - ☆ For each operation  $O_{ik}(x)$  of transaction  $T_i$ ,  $O_{ik}(x)$  is executed at site  $S^l$  where  $x \in D^l$ ,  $S^l$  sends an ack to  $S^i$
  - ☆ At the end of  $T_i$ , all participating sites terminate (a,c)  $T_i$



COMP-512: Distributed Systems

# Global and local Histories

- ❑ Distributed transaction processing produces local histories at each site
  - ☆ The local history at site  $S_j$  contains for each transaction  $T_i$  the operations that have been executed at  $S_j$ .
- ❑ The global history is the union of all local histories.
- ❑ A concurrency control mechanism guarantees global serializability if all global histories that might occur are conflict-serializable

COMP-512: Distributed Systems

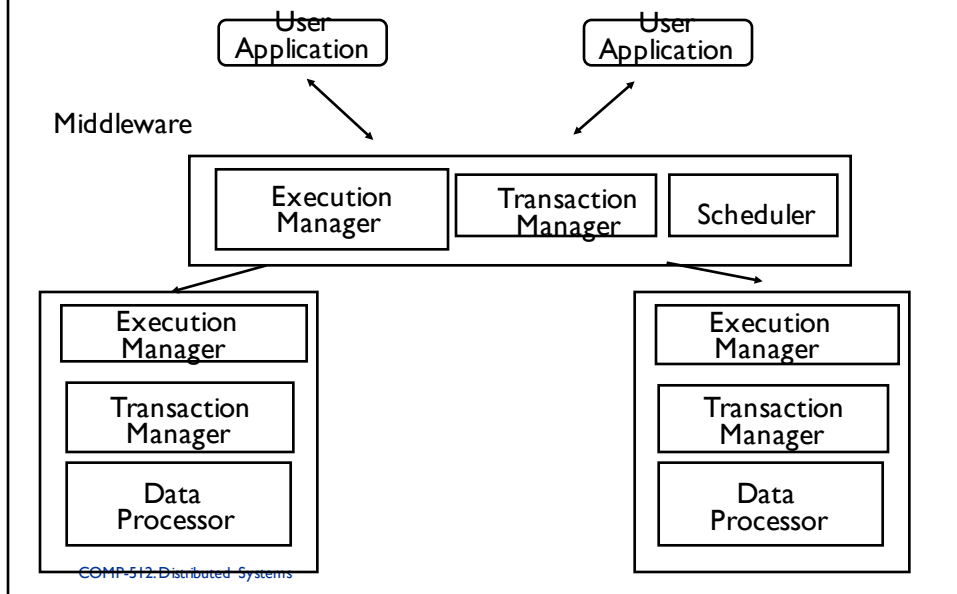
# Independent Execution

T1	T2	S1:x	S2:y	S3:z
r1(x)		r1(x)		
	r2(z)		r2(y)	r2(z)
w1(x)	r2(y)	w1(x)		
w1(y)		C1	w1(y)	
c1		w2(x)	c1	
	w2(x)	C2	c2	c2
	c2			

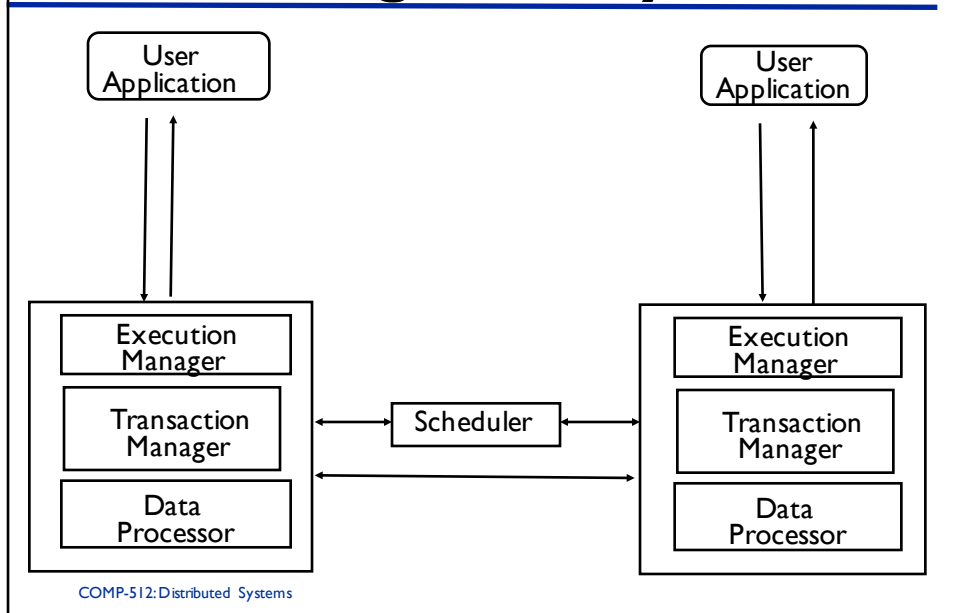
- ❑ Local histories are serializable, but the global history is not serializable

COMP-512: Distributed Systems

## Centralized Scheduler Middleware-based System

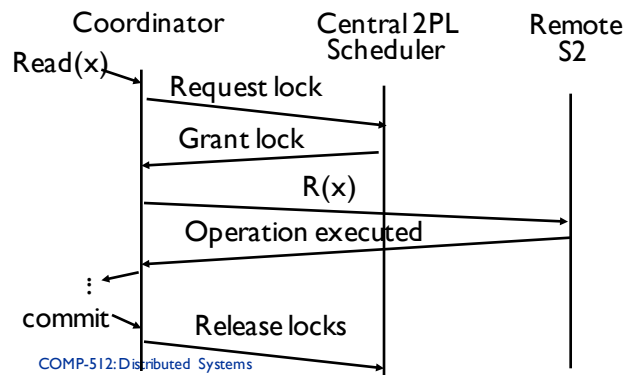


## Centralized Scheduler Homogenous System



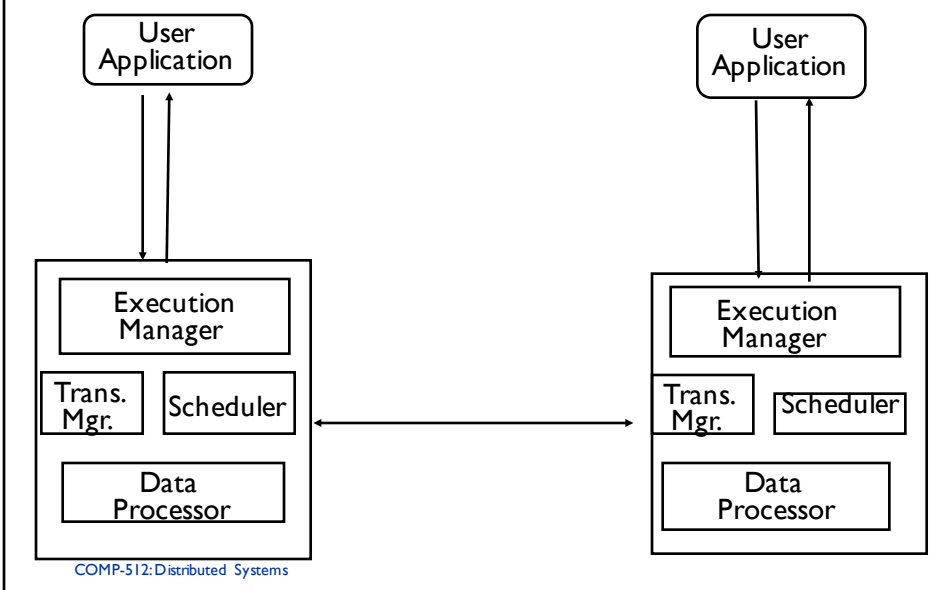
## Centralized 2PL

- There exists one global 2PL Scheduler
- Let  $S1$  be the server to which transaction  $T_i$  was submitted:
  - ☆ For each operation  $o_{ij}(X)$ ,  $T_i$  first requests the corresponding lock from the central 2PL scheduler
  - ☆ Once the lock is granted, the operation is forwarded to the server  $S2$  maintaining  $X$ .
  - ☆ In a middleware based system, global scheduler is typically part of the middleware



COMP-512: Distributed Systems

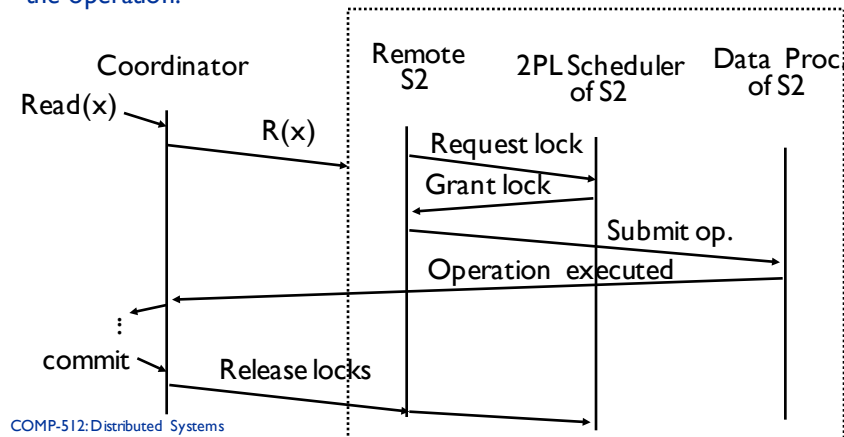
## Distributed Scheduler homogenous system



COMP-512: Distributed Systems

## Distributed 2PL

- ❑ Each server has its own local 2PL scheduler
- ❑ Let  $S_1$  be the server to which transaction  $T_i$  was submitted:
  - ☆ For each operation  $o_{ij}(X)$ ,  $TM_1$  forwards the operation to the  $TM_2$  of server  $S_2$  maintaining  $X$
  - ☆ The remote site first acquires a lock on  $X$  and then submits the execution of the operation.



COMP-512: Distributed Systems

## Deadlock in distributed system

- ❑ Timeout Mechanism
  - ☆ if a transaction waits for a lock longer than a predefined timeout interval, assume it is in a deadlock cycle and abort the transaction
  - ☆ Disadvantage: choice of adequate timeout interval is crucial
- ❑ Global Deadlock Detection
  - ☆ Every site sends its local graph periodically to a central site
  - ☆ Central site will eventually detect conflict
  - ☆ Suitable with centralized 2PL
- ❑ Deadlock Prevention
  - ☆ Lock all items at the beginning of transaction
- ❑ Distributed Deadlock detection
  - ☆ seldomly used

COMP-512: Distributed Systems

# Deadlock Detection (Continued)

Example:



T1: S(A), r(A),  
 T2: X(B), W(B)  
 T3: S(C), R(C)  
 T4: X(C)

S(B)

X(C)

**X(A)**

X(B)



COMP-512: Distributed Systems