

---

# Optimistic Concurrency Control

---

## Optimistic CC

---

- ❑ Locking is conservative
  - ☆ Locking overhead even if no conflicts
  - ☆ Deadlock detection/resolution (especially problematic in distributed environment)
- ❑ Optimistic concurrency control
  - ☆ Perform operation first
  - ☆ Check for conflicts only later (e.g., at commit time)
- ❑ Centralized systems never used textbook optimistic CC
  - ☆ More popular: snapshot isolation
- ❑ Interesting alternative for distributed environment


# Kung-Robinson Model

- ❑ Working Phase:
  - ☆ If first operation on X, then access the last committed version of X, make local copy and read/write local copy
  - ☆ Otherwise read/write read/write local copy
  - ☆ Keep WriteSet containing objects written
  - ☆ Keep ReadSet containing objects read
- ❑ Validation Phase
  - ☆ Check whether transaction conflicts with other transactions
- ❑ Update Phase
  - ☆ Upon successful validation, local copies of updated objects are made public. They are committed
- ❑ Assumption (for simplicity):
  - ☆ Only one transaction may be in validation phase and update phase at a time (e.g., implemented as critical section)

COMP-512: Distributed Systems

3

## Example



- ❑ Validation order determines serialization order
  - ☆ if serialization order not possible according to validation order then abort validation transaction
- ❑ Validation and update serially
  - ☆ All updates are automatically executed in order transactions are validated
  - ☆ Validation only needs to check write/read conflicts

COMP-512: Distributed Systems

4

# Backward Validation

- ❑ Maintain transaction counter TC, initialize TC := 0
- ❑ Upon begin of transaction  $T_i$  (could be first read/write request)
  - ☆  $StartTC(T_i) := TC$
  - ☆  $CommitTS(T_i) := \text{undefined}$
  - ☆  $WriteSet(T_i) := ReadSet(T_i) := \{\}$
- ❑ Upon  $r_i(x), w_i(x)$  request of  $T_i$ 
  - ☆ If first operation on X, access X (latest committed version) and make local copy
  - ☆ read/write local copy of X
  - ☆ If  $r_i(x)$ , then  $ReadSet(T_i) := ReadSet(T_i) \cup \{x\}$
  - ☆ If  $w_i(x)$ , then  $WriteSet(T_i) := WriteSet(T_i) \cup \{x\}$
- ❑ At time of validation of transaction  $T_i$  (in critical section)
  - ☆ For all  $StartTC(T_j) < j \leq TC$ 
    - Let WS be the write set of transaction T with  $CommitTS(T) = j$
    - If  $WS \cap ReadSet(T_i) \neq \{\}$ , then abort  $T_i$
  - ☆ If not aborted
    - $TC = TC + 1$
    - $CommitTS(T_i) = TC$
    - for each x in  $WriteSet(T_i)$  local copy of x becomes the official last committed version

# Discussion

- ❑ Validation and Update Phase in Critical section
  - ☆ Reduced concurrency
  - ☆ Optimizations possible
- ❑ Space overhead:
  - ☆ To validate  $T_i$ , must have WriteSets for all  $T_j$  where  $T_j < T_i$  and  $T_j$  was active when  $T_i$  began.
  - ☆ Each transaction must record read/write activity
- ❑ No deadlock but potentially more aborts.

# Snapshot Isolation

- ❑ Multiple versions
- ❑ Check write/write conflicts instead of read/writes
- ❑ Read snapshot as of start of transactions

# Snapshot Isolation

- ❑ Maintain transaction counter TC, initialize  $TC := 0$
- ❑ Upon begin of transaction  $T_i$  (could be first read/write request)
  - ☆  $StartTC(T_i) := TC$
  - ☆  $CommitTS(T_i) := \text{undefined}$
  - ☆  $WriteSet(T_i) := ReadSet(T_i) := \{\}$
- ❑ Upon  $r_i(x), w_i(x)$  request of  $T_i$ 
  - ☆ If first operation on  $X$ , access version with label  $TS(X)$  such that
    - $TS(x) \leq StartTC(T_i)$  and
    - No  $TS(x)$  with  $TS(x) < TS'(x) \leq StartTC(T_i)$
  - ☆ read/write local copy of  $X$
  - ☆ If  $w_i(x)$ , then  $WriteSet(T_i) := WriteSet(T_i) \cup \{x\}$
- ❑ At time of validation of transaction  $T_i$  (in critical section)
  - ☆ For all  $StartTC(T_j) < j \leq TC$ 
    - Let  $WS$  be the write set of transaction  $T$  with  $CommitTS(T) = j$
    - If  $WS \cap WriteSet(T_i) \neq \{\}$ , then abort  $T_i$
  - ☆ If not aborted
    - $TC = TC + 1$
    - $CommitTS(T_i) = TC$
    - for each  $x$  in  $WriteSet(T_i)$  write version  $X$  with label  $TC$  into the data store

## Snapshot Isolation vs. Classical Optimistic CC

---

- ❑ Do not need to keep track of reads
  - ☆ Typically many more reads than writes
  - ☆ Predicate reads difficult to grasp
- ❑ Read-only transactions never need validation
  - ☆ Are serialized at the time point they started
  - ☆ Readers never conflict with writers
- ❑ Natural for append-only stores
  - ☆ Have become popular in the recent past
- ❑ No serializability!!
- ❑ Oracle, PostgreSQL, Microsoft SQL Server, ...