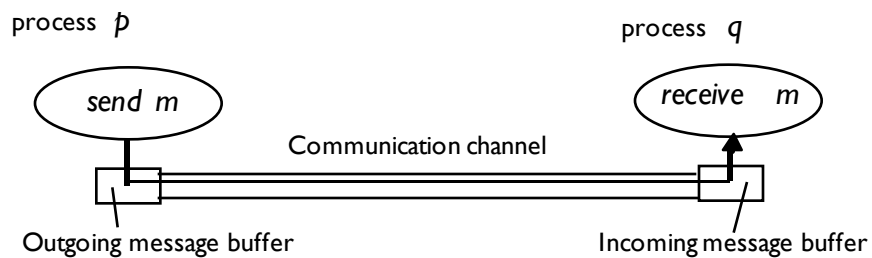


Communication Basics:

1) Inter-process Communication

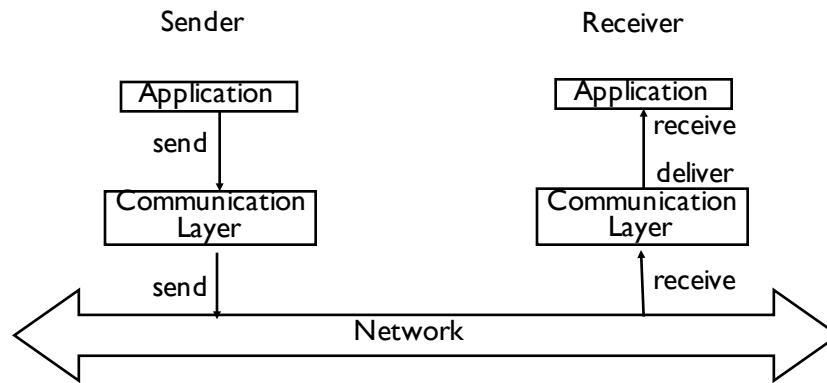
2) Internet Stack

Inter-Process Communication



- One communication “unit” consists of two primitives
 - ☆ The **send** primitive is called by the sending process (caller, sender)
 - ☆ A corresponding **receive** primitive must be called by the receiving process (callee, receiver)

Basic Architecture



COMP-512: Distributed Systems

3

Blocking vs. Non-blocking send

❑ Non-blocking send

- ☆ send operation returns and sending process is allowed to proceed as soon as the underlying communication layer has received the message and committed to process it

❑ Blocking send

- ☆ send operation only returns once it is confirmed that at the receiver site the message
 - was delivered to the receiving process or
 - was received by the receiving process or
 - will be delivered to the receiving process or
 - (depends on definition)

COMP-512: Distributed Systems

4

Blocking vs. Non-blocking receive

- ❑ Blocking receive
 - ☆ The receive primitive blocks until a message arrives
 - ☆ Multi-threaded environment:
 - one receiving listener thread that has loop with blocking receive
 - other threads do other work...
- ❑ Message Handler
 - ☆ the application process provides the communication layer with a handler routine
 - ☆ the communication layer calls this routing upon message delivery
- ❑ Note: Book in chapter 4 calls this synchronous/asynchronous communication; here in class we do NOT

Source and Destination Addresses

- ❑ Abstract: process P
- ❑ Internet
 - ☆ address (=host) + port (location dependent)
 - ☆ Port is a message destination within a computer,
- ❑ Object (location independent)
 - ☆ location of object is determined at runtime
- ❑ Service (location independent);
 - ☆ Service name is translated at runtime to server location

Performance

❑ Network Latency

- ☆ Delay between the start of transmission of a message over the network and the beginning of its receipt by receiving process
 - Time for the first bit to be transmitted through the network

❑ Process Latency

- ☆ Network Latency +
- ☆ Processing time at sender/receiver (communication layers)
 - Includes marshalling/unmarshalling
- ☆ Other delays due to load at sender/receiver, network

Performance

❑ Data transfer rate / bandwidth

- ☆ Speed at which data can be transferred between 2 computers
- ☆ Bits/sec

❑ Jitter:

- ☆ Variation in time taken to deliver a series of messages (often bursty delivery)

Synchronous vs. asynchronous

- ❑ Synchronous distributed system
 - ☆ each message is transmitted within a known bounded time
 - ☆ The time to execute each step of a process has known lower and upper bounds
 - ☆ Each process has a local clock whose drift rate from real time has a known bound
- ❑ Asynchronous distributed system
 - ☆ Message may need an arbitrary time to be transmitted
 - ☆ Each step of a process can take an arbitrary time
 - ☆ Clocks drift rates are arbitrary
- ❑ In the following: asynchronous model if not stated otherwise

Asynchronous Agreement

- ❑ Model:
 - ☆ Two armies A1 and A2 on top of two mountains. Enemy B in the valley
 - ☆ A1 and A2 can exchange messengers
 - synchronous: Each messenger needs at least *min* at most *max* time units
 - asynchronous but reliable: messengers will *eventually* arrive but there is no time limit on how long the messenger needs

Asynchronous Agreement

- Problem 1: A1 and A2 have to agree who starts attack

- ☆ Solution:

- both send message with number of soldiers,
- wait until receive the message of the other;
- the army with more soldiers starts attack

Asynchronous Agreement

- Problem 2: A1 and A2 have to agree **when** to attack; they have to attack nearly at the same time (difference at most X time-units)

Assume A1 has to start attack

- Asynchronous:

- ☆ A1 makes suggestion “at 2pm” or “in 20 minutes”.
- ☆ But no guarantee that message arrives before 2 pm or within 20 minutes.

- Synchronous model:

- ☆ A1 makes suggestion “attack now”, waits for *min* time units and attacks
- ☆ After receipt of message, A2 attacks
- ☆ A2 attacks at most (*max* - *min*) time units after A1
- ☆ If $(\text{max} - \text{min}) \leq X$, this works!

Possible failures

- ❑ **Omission failures:**
 - ☆ *process omission failure:*
 - *fail-stop:*
 - ▲ Process halts and remains halted;
 - ▲ Others can detect this state
 - *crash:*
 - ▲ Process halts and remains halted;
 - ▲ other processes may not be able to detect this
 - ☆ *communication omission failures*
 - receiving process does not receive message
 - typically because of
 - ▲ buffer overflow "somewhere": individual message loss
 - ▲ network partition
- ❑ **Arbitrary failures (byzantine failures)**
 - ☆ processes might not follow agreed protocol
 - ☆ messages might be corrupted
 - can be solved by checksums
- ❑ **Process failures vs. network partitions:** In general, it is impossible to distinguish between a process failure and a communication failure.

COMP-512: Distributed Systems

13


Reliability

- ❑ **Reliable delivery:**
 - ☆ **Validity:**
 - any message sent is eventually delivered to the application (unless the application fails)
 - ☆ **Integrity:**
 - the message received is identical to the one sent, and no messages are delivered twice
- ❑ **Non-reliable delivery:**
 - ☆ The communication primitive only provides a best effort. That is, when no failures occur, the message will arrive. But in case of failures, it might not arrive, be duplicated, etc
 - ☆ message identity typically provided
- ❑ **Careful !!!**
 - ☆ When you design a communication protocol with reliable delivery you must specify the failures that the protocol can handle
 - ☆ The protocol must work correctly if the specified failures occur
 - ☆ The protocol may not work if unspecified failures occur

COMP-512: Distributed Systems

14

Layered Architecture



COMP-512:Distributed Systems

15

Simple Example

□ Assume that

- ☆ Processes never fail
- ☆ No network partitions
- ☆ Network provides unreliable $n_send(q, message) / n_receive(message)$
 - Messages can get lost

□ Reliable communication module

- ☆ Implement reliable send/receive using above unreliable $n_send/n_receive$ pair
- ☆ Basic ideas
 - Acknowledgements
 - Resubmissions
 - Message Identifiers


COMP-512:Distributed Systems

16

Extended

- ❑ What if processes can crash and messages can get lost?

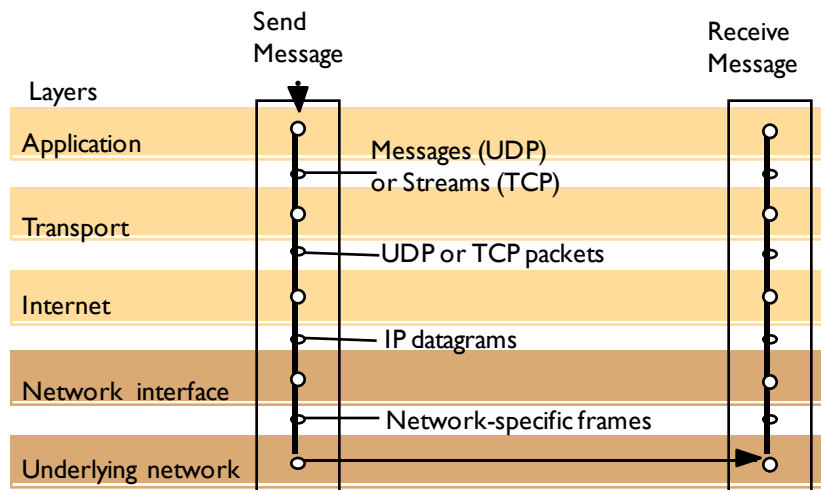
Networks



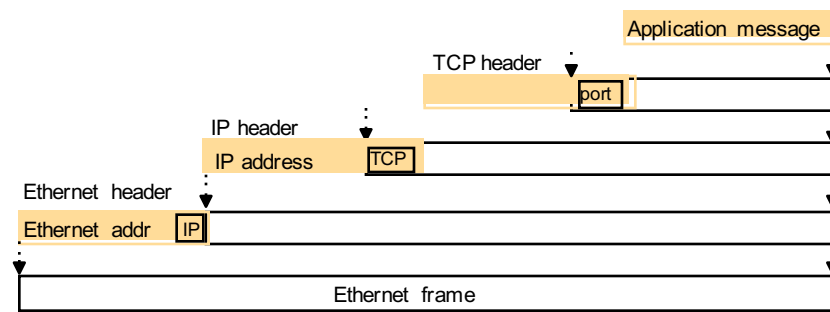
Network Types

	Range	Bandwidth/ Throughput (Mbps)	Latency (ms)
LAN	1-2 kms	10-10000	0.1-10
WAN	worldwide	0.010-600	100-500
MAN	2-50 kms	1-600	10
Wireless LAN	0.15-1.5 km	11-300	5-20
Wireless WAN	worldwide	348-14.4	100-500
Internet	worldwide	0.5-600	100-500

Communication Layers: Example TCP/IP



Message Format



Addresses

- ❑ IP Addresses: 32 bits
 - ☆ Network part: identifies subnet
 - ☆ Host part: identifies host on subnet
 - ☆ Hierarchical assignment
 - Organization is given a subnet identifier
 - Organization assigns host identifiers within subnet
- ❑ Routers/Bridges connect different subnets
 - ☆ they have an IP address for each subnet
- ❑ Addressing for portable computers
 - ☆ NAT: network address translation
 - ☆ unregistered addresses (not unique throughout the network)
 - ☆ special router does mapping
- ❑ Network addresses (physical addresses)


Network layer: Ethernet

- ❑ send message packet from node A to node B in same LAN
 - ☆ broadcast message on medium
 - ☆ controller hardware at each node listens on medium
 - ☆ pick up messages that are addressed to them and gives them
 - ☆ header: sender/destination (network addresses), ...
- ❑ Packet Collision possible
 - ☆ resubmit after waiting certain time

IP protocol

- ❑ host-to-host transmission
- ❑ uses IP address
- ❑ unreliable, best-effort
 - ☆ Message can get lost somewhere
- ❑ Tasks:
 - ☆ transform IP address to network address of underlying network protocol
 - ☆ ROUTING

Routing in a wide area network



COMP-512:Distributed Systems

25

Routing Table

Routings from A			Routings from B			Routings from C		
To	Link	Cost	To	Link	Cost	To	Link	Cost
A	local	0	A	1	1	A	2	2
B	1	1	B	local	0	B	2	1
C	1	2	C	2	1	C	local	0
D	3	1	D	1	2	D	5	2
E	1	2	E	4	1	E	5	1

Routings from D			Routings from E		
To	Link	Cost	To	Link	Cost
A	3	1	A	4	2
B	3	2	B	4	1
C	6	2	C	5	1
D	local	0	D	6	1
E	6	1	E	local	0

COMP-512:Distributed Systems

26

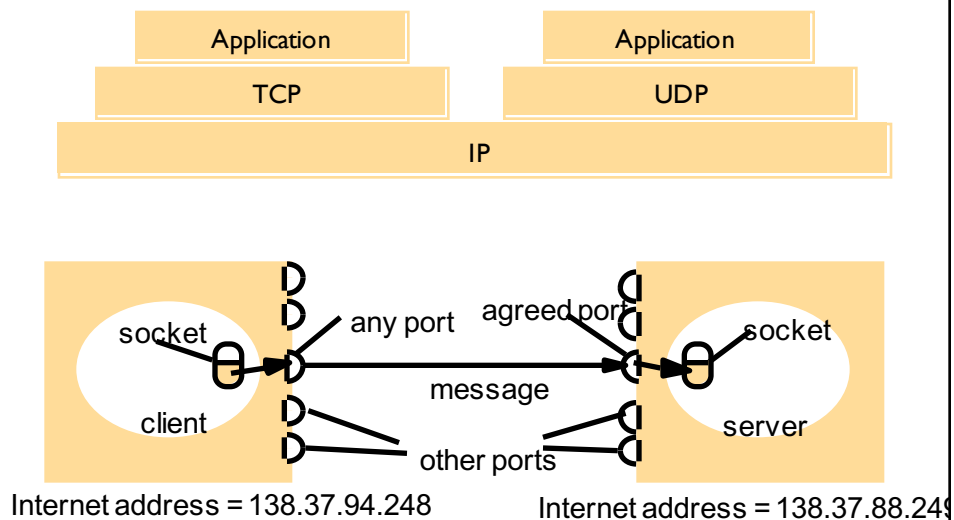
Issues

- ❑ Fast lookup of next hub
- ❑ maintenance
 - ☆ if link broken, set cost to infinity
 - ☆ send periodically routing table to others
 - others update their tables
- ❑ reasonable size
 - ☆ geographically oriented IP addresses
 - ☆ defaults

COMP-512: Distributed Systems

27


The programmers view



COMP-512: Distributed Systems

28

The programmers view



□ UDP

- ☆ thin software layer on top of IP which builds interface to application
- ☆ same functionality as IP: unreliable, best-effort, message-oriented
- ☆ checksums against message corruption: integrity
- ☆ domain name resolution: application can use logical name (mimi.cs.mcgill.ca) instead of IP address
 - more on how domain names are translated into IP addresses later
 - ▲ (it's a distributed application by itself!)
- ☆ connection-less

UDP

- send(message byte string, message-length, IP receiver, port receiver)
 - ☆ UDP then adds IP sender, port sender
- receive(byte string into buffer, length, IP sender, IP port)
- Java
 - ☆ DatagramPacket: message, message length, IP address, port
 - IP address and port have different content for sending / receiving
 - ☆ DatagramSocket:
 - constructor with port as input (for server)
 - constructor without port as input (for client)
 - send and receive methods with DatagramPacket as argument

TCP

- ❑ connection-oriented:
 - ☆ Both processes first run a connection protocol
 - Handshake that both agree on building a communication channel
 - ☆ Only after connection is established, application level messages can be sent in both directions
- ❑ stream-oriented
 - ☆ both parties have input and output buffer
- ❑ reliable:
 - ☆ stream split into sequence of data segments (TCP messages) with sequence numbers
 - ☆ acknowledgement scheme:
 - receiver sends to sender periodically
 - ▲ highest sequence number in input stream + window size
 - Flow Control to avoid buffer overflow at receiver
 - ▲ sender can only send window size of data before next acknowledgement
 - retransmission of lost messages
 - ▲ sender keeps messages in output buffer until acknowledgment receives from receiver
 - ▲ if message not acknowledged within specified timeout, then it retransmits it.
 - ▲ duplicate detection using sequence numbers
- ❑ FIFO delivery
 - ☆ sequence numbers allow detection of out-of-order messages and reorder them before delivering them to the application

COMP-512: Distributed Systems

31

TCP

- ❑ Overview
 - ☆ Client creates socket with input and output stream
 - ☆ Server has special socket object with associated port
 - Server listens on socket for connection request message from client
 - Clients makes connection request to this socket
 - upon receiving request from Client C
 - ▲ create new socket with input and output stream only for communication with client C
- ❑ Java
 - ☆ ServerSocket: constructor has port as input
 - accept method returns new instance of Socket class which is already connected to client socket
 - ☆ Socket:
 - one constructor (used by client) has as input name and port of server
 - ▲ creates socket with local port and connects to ServerSocket
 - has associated input and output streams of certain types
 - ▲ (e.g., DataInputStream, ObjectInputStream, ...)
 - input/output streams have methods: write/read

COMP-512: Distributed Systems

32

Data Representation in Messages

- ❑ Messages are a sequence of bytes -> data must be flattened
- ❑ Marshalling/serialization: transformation of a collection of data items into a form suitable for transmission
- ❑ Unmarshalling: transforming a message back into a collection of data items
- ❑ Common Mechanisms:
 - ❑ Java Serialization
 - ❑ XML / JSOC
 - ❑ CDR / Corba
 - ❑ Google Protocol Buffer

COMP-512:Distributed Systems

33

Java Declaration

```
Class Person implements Serializable {  
    int id;  
    string name;  
    string place;  
    int year;  
  
    Public Person (int pid, string pname, string pplace, int pyear) {  
        id = pid;  
        name = pname;  
        place = pplace;  
        year = pyear;  
    }  
  
    Person myPerson = new Person(12345, "Smith", "London", 1934);  
}
```

COMP-512:Distributed Systems

34

Java usage

❑ To exchange object of class Person via TCP

☆ At sender

☆ `ObjectOutputStream oos = new ObjectOutputStream(socket);`

☆ `oos.writeObject(myPerson);`

☆ At receiver


☆ `ObjectInputStream ois = new ObjectInputStream(socket);`

☆ `readPerson = (Person) ois.readObject();`

JSON

❑ textual format to present structured data

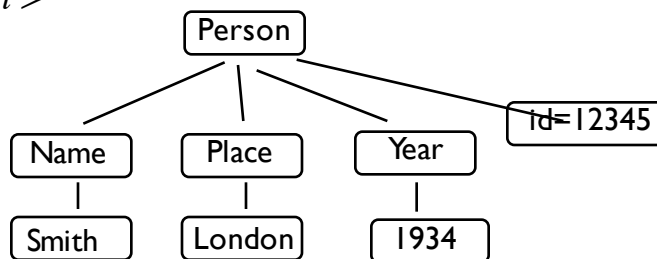
```
{  
  "person" : {  
    "id" : 12345  
    "name" : "Smith"  
    "place" : "London"  
    "year" : 1934  
  }  
}
```



XML

- ❑ textual format to present structured data

```
<person id="12345">  
  <name>Smith</name>  
  <place>London</place>  
  <year>1934</year>  
  <!-- a comment -->  
</person>
```



COMP-512: Distributed Systems

37

XML

- ❑ elements
 - ☆ basic building blocks
 - ☆ have start and end tag which indicate the name of element
 - ☆ elements can be nested
- ❑ attributes
 - ☆ element can have attributes
 - ☆ attribute is name/value pair
 - ☆ embedded in start tag of element
- ❑ text
 - ☆ part of an element
 - ☆ represents the content

COMP-512: Distributed Systems

38

Json/XML

- ❑ There exist many parsers for JSON/XML
 - ❑ e.g., Java software for conversion between Java objects and XML/JSON .

Google Protocol Buffer

```
message Person {  
    required int32 id = 1;  
    required string name = 2;  
    optional string location = 3;  
    optional int place = 4;  
}  
Person person;  
person.set_name("John Doe");  
person.set_id(1234);  
person.set_location("London");
```

Google Protocol Buffer

- ❑ Programming Similar to Java

- ☆ Output

```
fstream output("myfile", ios::out | ios::binary);  
person.SerializeToOstream(&output);
```

- ☆ Input

```
fstream input("myfile", ios::in | ios::binary);  
Person person;  
person.ParseFromIstream(&input);
```

XML/JSON

- ❑ Transmitted in Ascii / Text Format
- ❑ Contains meta data and data

Internal Representation

❑ XML/JSON

- ☆ Test / ascii
- ☆ Human readable
- ☆ Contains meta data and data

❑ Java

- ☆ Serialized binary format
- ☆ Contains information about data types

Serialized values

Explanation

Person	8-byte version number		h0
3	int year	java.lang.String name:	java.lang.String place:
1934	5 Smith	6 London	h1

class name, version number

*number, type and name of
instance variables*

values of instance variables

The true serialized form contains additional type markers; h0 and h1 are handles

COMP-512: Distributed Systems

43

Internal Representation

❑ Google Protocol Buffer

```
message Person {
    required int32 id = 1;
    required string name = 2;
    optional string location = 3;
    optional int32 year = 4;
}

Person person;
person.set_id(1234);
person.set_name("Smith");
person.set_year(1934);
```

❑ Internal binary presentation

- ☆ Contains coded info about (1) data type (2) attribute (3) real data

0 1 1234

2 2 "Smith"

0 4 1934

COMP-512: Distributed Systems

44

CORBA's Data Representation

<i>index in sequence of bytes</i>	<i>← 4 bytes →</i>	<i>notes on representation</i>
0–3	5	<i>length of string</i>
4–7	"Smit"	<i>'Smith'</i>
8–11	"h _ _ _"	
12–15	6	<i>length of string</i>
16–19	"Lond"	<i>'London'</i>
20–23	"on _ _"	
24–27	1934	<i>unsigned long</i>

The flattened form represents a *Person* struct with value: {'Smith', 'London', 1934}

CORBA Marshalling

- ❑ message contains only the data but not the information about data types
- ❑ programmer can specify the structs that are the content of messages in an IDL (interface definition language) file
 - ☆ CORBA generates automatically marshalling and unmarshalling operations for these structs