


---

**Request Reply**  
**Remote Method Invocation**  
**Remote Service Invocation**

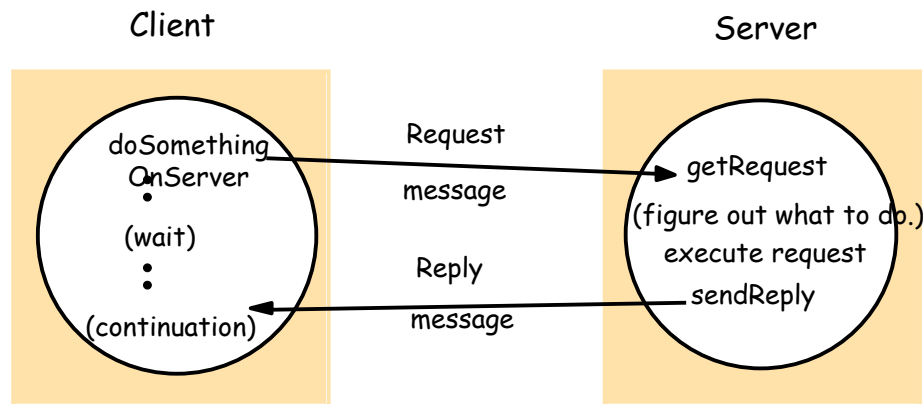
---

## **Middleware layers**

---



# Client-Server communication



512: Distributed Systems

# Implementation of RR over UDP

- ❑ Failure model:
  - ☆ *message loss, process crash*
- ❑ Invocation semantics:
  - ☆ *maybe*
    - *provides best-effort*
    - *avoids blocking by timeout*
      - ▲ *give client failure exception after timeout*

512: Distributed Systems

# Implementation of RR over UDP

## □ Invocation semantics: at-least-once

- ☆ client software
  - resubmits request if it does not receive result within timeout interval (up to N submissions)
  - returns after receiving first reply (ignores further replies it might get)
- ☆ Server might execute request more than once
- ☆ acceptable if service is *idempotent*
  
- ☆ client resubmits up to N times and then gives up
  - Up to N-1 message losses are accepted by the protocol
  - Client cannot distinguish whether request or reply messages are lost or whether server has crashed
  - After N timeouts it does not know whether server has executed request or not

512: Distributed Systems

# Implementation of RR over UDP

## □ Invocation semantics: at-most-once (or better exactly-once)

- ☆ client software
  - Same as at-least-once
- ☆ server must detect duplicate requests
  - unique request identifiers
    - ▲ typical: IP address of client + process id + sequence id within process
  - server must keep all requests (*history*)
  - once reply is generated, reply is added to corresponding request
  - upon receiving request again,
    - ▲ reply is returned without reexecution
    - ▲ note: if timeout too small, server might receive second submission of request while executing first: must detect this!
  
- ☆ garbage collection at server?
  - Client sends ack when it has received response (RRA)
  - Similar optimizations as TCP

512: Distributed Systems


# Comparison

- ❑ Compare at-most-once RR over UDP with TCP
  - ☆ Similar mechanisms but slightly different implementation
    - Exactly-once service execution vs. exactly-once message delivery
- ❑ Compare at-most-once RR over UDP with RR using simply TCP
  - ☆ Both provide the same semantics (exactly-once as long as N tries are sufficient)
  - ☆ Performance different
    - Different number of messages sent

512: Distributed Systems

# HTTP: Hypertext Transfer Protocol

- ❑ On top of TCP



512: Distributed Systems

# Features

---

- ❑ Request is on a Resource
- ❑ Uniform Resource Identifier / Uniform Resource Locator (URL)
  - ☆ Location
    - Logical address name plus potentially port
  - ☆ Unique Name within location
    - Can be a file (static content)
    - Can be a program
  - ☆ <http://www.cs.mcgill.ca/~kemme/cs512/index.html>
- ❑ Most common:
  - ☆ Get Request to retrieve static content

512: Distributed Systems


# Advanced Requests

---

- ❑ Methods on the resource
  - ☆ Get
    - On data: return content
    - On program: run program and return result
  - ☆ Put
    - URL already exists: override
    - New URL: create
  - ☆ Post
    - Add data to the resource (e.g. a mailing list, database...)
  - ☆ Others
    - ☆ Parameters: program parameters
      - E.g., data entered in a form...
      - Syntax can get complicated!
- ❑ Textual Format

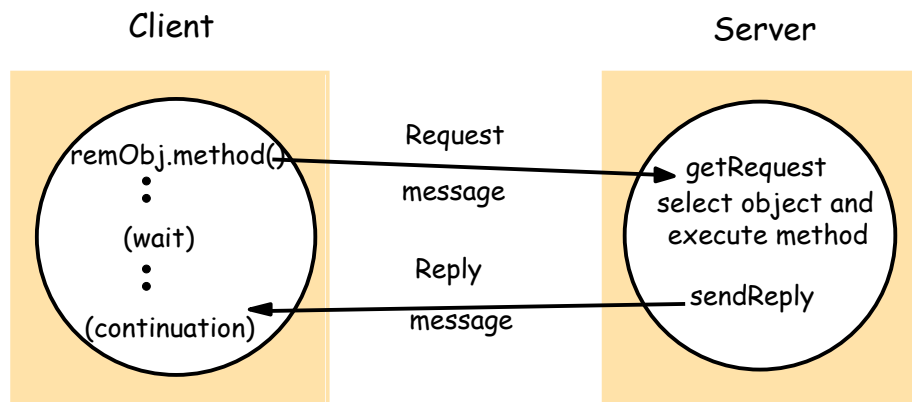
512: Distributed Systems

# Remote Method Invocation



512: Distributed Systems

# Client-Server communication



512: Distributed Systems

# Programming Models for Distributed Applications

- ❑ Remote Procedure Call (RPC)
  - ☆ Client calls a procedure in a server running on a different machine
  - ☆ Idea: call to remote procedure should look like a call to local procedure
  - ☆ Theory since 1976, first implementations in the early 80's
- ❑ Remote Method Invocation (RMI)
  - ☆ Same as RPC in OO-world
  - ☆ Object calls method of another object residing on a different machine
  - ☆ First: implementations had same architecture than RPC
  - ☆ Current developments: more features and variations
  - ☆ Trend to truly distributed computing

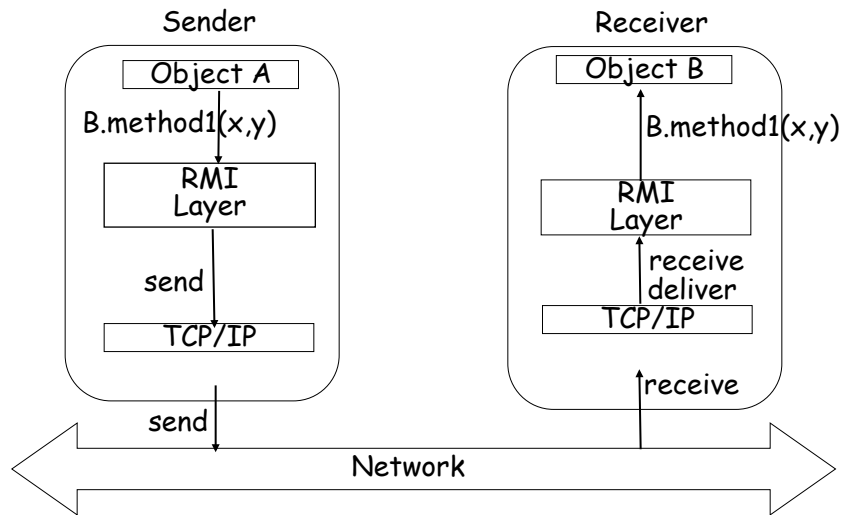
512: Distributed Systems

## Review of OO-programming

- ❑ An object encapsulates both *data* and *methods*
- ❑ Objects are accessed via *object references*
- ❑ Method Invocation can cause
  - ☆ the state of the called object to change
  - ☆ further invocations on methods in same or other objects
- ❑ Access to variables (direct/indirect)
  - ☆ direct (e.g., public): not suitable in distributed world
  - ☆ indirect (e.g., via getter/setter methods): encapsulation
  - ☆ Assumption from now: variables of an object can only be accessed via method calls
- ❑ *Exceptions*
  - ☆ are thrown when an error occurs.
  - ☆ can be caught (redirection to specific code handling the exception)
  - ☆ or delivered to caller.
- ❑ *Interfaces*:
  - ☆ provide signatures of a set of methods
  - ☆ description of input and output parameters of the methods
- ❑ *Garbage collection*
  - ☆ frees space when objects are no longer needed

512: Distributed Systems

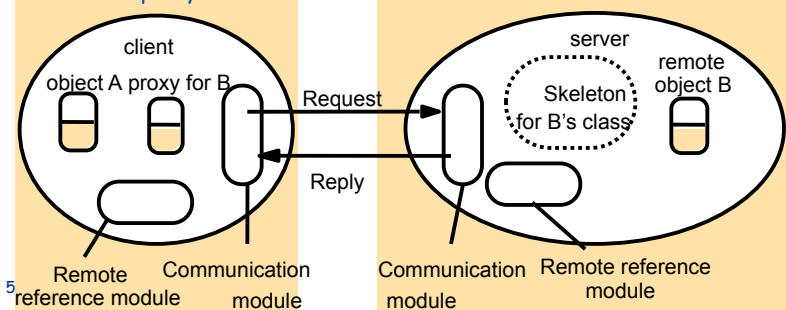
# RMI I



512: Distributed Systems

# RMI I

- ❑ *Remote Interface File* of an object class
  - ☆ Interface file for the set of methods that can be called remotely
  - ☆ object class can have further methods that can only be called locally
- ❑ Method description in remote interface file
  - ☆ input and output parameters are specified
  - ☆ a parameter can be an object
- ❑ Implementation of local and remote methods
  - ☆ only at server
- ❑ Proxy object at client (stub)
  - ☆ contains proxy methods for all methods of the remote interface





## RMI II

### □ On the client

- ☆ object A calls method of remote object B
  - it internally calls the proxy-method of the local proxy object.
  - proxy also called *stub*
  - i.e., object A has a object reference to the local proxy object
- ☆ Within proxy-method and underlying communication module (RR module)
  - marshals *request* message containing
    - ▲ *remote reference* to the remote target object,
    - ▲ its own methodID
    - ▲ and the arguments
  - sends request message to server
  - awaits reply message
  - unmarshals reply
  - returns the results to the invoker

512: Distributed Systems

## RMI III

### □ On the server

- ☆ RMI communication module receives message
- ☆ Determines skeleton of object B given the reference and calls skeleton's method based on methodID
- ☆ Skeleton method unmarshals message and invokes corresponding method on real object B
- ☆ marshals the result (together with any exception) into result message
- ☆ Communicatin module sends result message back to RMI layer on caller
  
- ☆ Proxies and skeletons are automatically generated with special compiler

512: Distributed Systems


# Interface File in Java RMI

```
import java.rmi.Remote;  
import java.rmi.RemoteException;  
public interface NumberStack extends Remote {  
    public int getNumber() throws RemoteException;  
    public void putNumber(int num) throws RemoteException; }
```

- ❑ By extending `java.rmi.Remote`, this interface's methods can be called from any virtual machine. Any object that implements this interface becomes a remote object.
- ❑ As a member of a remote interface, the `getNumber` method is a remote method. Therefore the method must be defined as being capable of throwing a `java.rmi.RemoteException`.
- ❑ This exception is thrown by the RMI system during a remote method when communication failure or a protocol error has occurred.
- ❑ Any code making a call to a remote method needs to handle this exception by either catching it or declaring it in its throws clause.

512: Distributed Systems

## A remote object and its remote interface



512: Distributed Systems

# Interface Definition Languages

- ❑ Java has the concept of interfaces and hence, the remote interface can be described as usual + extending interface to be REMOTE
- ❑ In general interfaces are described with Interface Definition Languages
- ❑ Examples RMI
  - ☆ Corba IDL
  - ☆ DCOM IDL for Microsoft's Distributed Common Object Model (DCOM) RMI
- ❑ Examples RPC
  - ☆ Sun XDR
  - ☆ DCE IDL: RPC system of OSF's Distributed Computing Environment DCE

512: Distributed Systems

# Remote Object Reference

32 bits	32 bits	32 bits	32 bits
Internet address	port number	time	objectID

- ❑ In non-distributed execution, objects can be accessed via object reference.
- ❑ In Java, a variable that appears to hold an object actually holds a reference to that object.
- ❑ An object A can invoke the method of a remote object B, if it has access to B's remote object reference.
- ❑ Generally: A remote object reference is an identifier that can be used throughout a distributed system to refer to a particular unique remote object.
- ❑ Note: In some RMI online information the remote object reference actually refers to a stub/proxy instance
- ❑ Conceptually like a URL

512: Distributed Systems

# Remote reference module

## ❑ Maintains remote object table

### ☆ entry


- remote object reference and local proxy reference (on client)
- remote object reference and real object reference (on server)

## ❑ Some Tasks

- ☆ On a server: when a remote object reference arrives in a request message it looks for the corresponding local object reference which points to the real object
- ☆ On a client: when a remote object reference arrives in a reply message it looks for the proxy object.
- ☆ When a remote object is to be passed as argument or result for the first time, the remote reference module creates a remote object reference and adds it to the table

512: Distributed Systems

# A basic Java RMI call



512: Distributed Systems

## The binder: RMI Registry

---

- ❑ An instance of the Rmiregistry must run on every computer on which a server wants to export a remote interface.
- ❑ It maintains a table mapping textual references to remote objects hosted on that computer.
  - ☆ `//computeName:port/objectName`
  - ☆ `computeName` is the host name on which the registry resides and `port` is the port the `rmiregistry` listens on

512: Distributed Systems

## RMI Registry Interface

---

- ❑ `Void rebind(String name, Remote obj)`
  - ☆ Register object by name
  - ☆ Override previous registration
- ❑ `Void bind (String name, Remote obj)`
  - ☆ Register an object by name
  - ☆ If existent throw exception
- ❑ `Void unbind(String name, remote obj)`
- ❑ `Remote lookup(String name)`
  - ☆ Used by client
  - ☆ Remote object reference is returned
- ❑ `String[] list()`
  - ☆ Show all names bound in this registry

512: Distributed Systems

# Method Parameters

- ❑ Input and output parameters can be
  - ☆ Primitive types (pass by value)
  - ☆ Remote objects (pass the remote object reference)
    - If the recipient of a remote object reference does not have the object proxy, then the object proxy is automatically downloaded
  - ☆ Local objects. They are passed by copy, using object serialization (the object must implement the `java.io.Serializable` Interface).
    - If the recipient of an object passed by value does not already possess the class of the object, its code is automatically downloaded.

512: Distributed Systems

# Distributed Garbage Collection

- ❑ An object must exist as long as there is a local or remote reference held
- ❑ Whenever the last reference is deleted the object should be garbage collected.
- ❑ A classical distributed algorithm problem:
  - ☆ Solution: reference counting
- ❑ 2-level
  - ☆ local proxy garbage collection on each client
    - Whenever a remote object reference enters a process, a proxy is created, when the last reference to this object on the process is deleted, the proxy is deleted
  - ☆ object maintenance on the server
    - whenever there is no proxy anymore and there is no local reference, you can delete the object

512: Distributed Systems

# Distributed Garbage Collection

- ❑ Algorithm outline:
  - ☆ Each server process maintains a list of processes that currently hold a remote object reference
  - ☆ Whenever a new proxy is created (a client first requests a reference), the remote reference module on the client first calls a function `addRef` on the server and then creates the proxy
  - ☆ When a proxy on a client is no more needed (detected by the local GC of the client), the remote reference module calls `removeRef` on the server
  - ☆ When the list is empty and there are no local references the remote object is removed
- ❑ Pair-wise request-reply between remote reference modules
- ❑ Only called when proxies are created/deleted
- ❑ Fault-tolerance has to be addressed
  - ☆ Idempotent `addRef` and `removeRef`
  - ☆ `removeRef` is correct whether `addRef` worked or not
  - ☆ Leases (max. time to live) in case `removeRef` gets lost

512: Distributed Systems

# Web Service: Interface

- ❑ Remote *Service Call*
  - ☆ Set of services / procedures / methods that can be called
- ❑ Service Description (all XML)
  - ☆ WSDL: web-service definition language
  - ☆ Service
    - Described by the format of request / reply messages
  - ☆ Bindings
    - Details about protocols
  - ☆ Endpoint address
    - Different to RMI!!

512: Distributed Systems

## Web-Service: Communication

---

- ❑ SOAP: Simple Object Access Protocol
- ❑ Messages in XML
  - ☆ SOAP messages have a specific format
- ❑ Protocol Pattern
  - ☆ Request/reply
    - Based on HTTP
    - HTTP contains endpoint address and action for efficiency
  - ☆ Other models are offered
    - In-only
    - Out-only
    - Reliable versions

512: Distributed Systems

## Web-Service with Java

---


- ❑ Java API for developing web-services: JAX-RPC
- ❑ Hides all the details of web-services
  - ☆ Define your interface using Java interface
    - Some restrictions for parameter types (objects are allowed, but no remote objects)
  - ☆ Web-service is a single object that offers the services

512: Distributed Systems



# Web-Service: Components

- ❑ Communication Module: Http
- ❑ Container:
  - ☆ Dispatcher calls skeleton (according to http header URL)
  - ☆ Skeleton (per servlet) translates SOAP message to java servlet call; translates reply back to SOAP



512: Distributed Systems