

Time

Hardware and Software Clocks

- ❑ At real time, t , the OS reads the time on the computer's hardware clock $H_i(t)$
- ❑ It calculates the time on its software clock
$$C_i(t) = \alpha H_i(t) + \beta$$
 - ☆ E.g. a 64 bit number giving nanoseconds since some base time
- ❑ Clock resolution
 - ☆ How accurate does it need to be?
- ❑ Clock Drift: 1 sec in 11 days
- ❑ Monotonicity
 - ☆ $t' > t \Rightarrow C(t') > C(t)$
 - ☆ can achieve monotonicity with a hardware clock that runs fast by adjusting the values of α and β in $C_i(t) = \alpha H_i(t) + \beta$


Synchronization

- ❑ Internal synchronization within bound D
 - ☆ given a set of computer clocks $C_i(t)$, $i = 1, 2, \dots, N$
 - ☆ $|C_i(t) - C_j(t)| < D$ for $i = 1, 2, \dots, N$
- ❑ Coordinated Universal Time (UTC)
 - ☆ based on very accurate physical clocks
 - ☆ Computer need receivers;
 - 0.1-10 milliseconds accurate for land-based reception
 - 1 microsecond through GPS
- ❑ External synchronization within bound D
 - ☆ Given a source S of UTC time and computer clock C
 - ☆ $|S(t) - C(t)| < D$

COMP 512: Distributed Systems

Synchronizing clocks in an asynchronous system


- ❑ Ideas:
 - ☆ use a time server S that receives signals from a UTC source
- ❑ Synchronizing process p with time server
 - ☆ Process p requests time in m_r and receives t in m_t from S
 - ☆ p sets its clock to
 - $t + T_{\text{round}}/2$
- ❑ Build average over several requests



COMP 512: Distributed Systems

Network Time ProtocolNTP

- ❑ Primary Servers are connected to UTC sources
- ❑ Secondary servers are synchronized to primary servers
- ❑ lowest level servers are user's computers
- ❑ accuracy of
 - ☆ 10s of milliseconds for Internet
 - ☆ 200 microseconds in Intranet



COMP 512: Distributed Systems


Events and Logical time

- ❑ A distributed system consists of a collection of N processes $p_i, i=1, 2, 3, \dots, N$
- ❑ Each process executes on a single processor.
- ❑ Information exchange only occurs through message exchange
- ❑ Each process p_i has a state s_i . The state changes according to the operations the process performs. (the state i.e., is the content of all application variables).
- ❑ Each process executes a sequence of events/actions,
 - ☆ send, receive or change own state
- ❑ Ordering of events in a distributed system
 - ☆ The sequence of events within a single process can be placed in a total order:
 - $e \rightarrow_i e'$ if and only if event e occurs before e' at p_i .
 - Whenever a message is sent, the event of sending a message occurs before the event of receiving the message
- ❑ A history of process p_i is a series of events ordered by \rightarrow_i

$$history(p_i) = h_i = \langle e_i^0, e_i^1, e_i^2, \dots \rangle$$

COMP 512: Distributed Systems


Events in a distributed system



COMP 512: Distributed Systems

Logical time

- We define the happened-before relation, denoted by \rightarrow , as follows
 - ☆ If there is a process p_i : $e \rightarrow_i e'$, then $e \rightarrow e'$
 - ☆ For any message m , $\text{send}(m) \rightarrow \text{receive}(m)$ (where $\text{send}(m)$ is the event of sending the message and $\text{receive}(m)$ the receiving event)
 - ☆ If e , e' , and e'' are events such that $e \rightarrow e'$ and $e' \rightarrow e''$, then $e \rightarrow e''$
 - ☆ If neither $e \rightarrow e'$ nor $e' \rightarrow e$, then e and e' are concurrent $e \parallel e'$



COMP 512: Distributed Systems

Logical clocks (Lamport clocks)

- ❑ Each process p_i maintains a logical clock L_i .
- ❑ L_i is a monotonically increasing counter (no relationship to real time)
- ❑ Each site timestamps events with the value of its logical clock.
- ❑ We denote the timestamp of event e at p_i by $L_i(e)$ (and omit the index if we speak about any event in the system)
- ❑ Goal: whenever $e \rightarrow e'$, then $L(e) < L(e')$, i.e., when e happens before e' , then e should have a smaller timestamp than e' .

COMP 512: Distributed Systems


Logical clocks (Lamport clocks)

- ❑ Algorithm at process p_i when event e takes place:

```
 $L_i := L_i + 1$ 
If Local event  $e$ :
     $L_i(e) := L_i$ 
Else if Send( $m$ ) event  $e$ :
     $L_i(e) := L_i$ ;
    Piggyback  $L_i(e)$  on  $m$  ( $m, L_i(e)$ )
Else if Receive( $m, L_j(e')$ ) event  $e$ :
     $L_i := \max(L_i, L_j(e') + 1)$ ;
     $L_i(e) := L_i$ 
```

COMP 512: Distributed Systems

Events in a distributed system



- It is easy to see that
 - ☆ If $e \rightarrow e'$ then $L(e) < L(e')$
- The opposite does not hold
 - ☆ If $L(e) < L(e')$, then this does not mean that $e \rightarrow e'$, instead it could also be that $e \parallel e'$

COMP 512: Distributed Systems

Vector Clocks

- Goal: find a timestamp V such that
 - ☆ $e \rightarrow e' \Leftrightarrow V(e) < V(e')$
 - ☆ In particular, this means, that if neither $e \rightarrow e'$ nor $e' \rightarrow e$, then neither $V(e) < V(e')$ nor $V(e') < V(e)$ (the timestamps are incomparable)
- Each processor p_i has a vector clock V_i which is an array of N integers. (one for each processor)
- We compare vector timestamps as follows
 - ☆ $V = V'$ iff $V[j] = V'[j]$ for $j = 1, 2, \dots, N$
 - ☆ $V \leq V'$ iff $V[j] \leq V'[j]$ for $j = 1, 2, \dots, N$
 - ☆ $V < V'$ iff $V \leq V'$ and not $V = V'$

COMP 512: Distributed Systems

Vector Clocks

□ Algorithm at process p_i

Upon start: $V_i[j] := 0$, for $j = 1, 2, \dots, N$

Upon event e :

$V_i[i] := V_i[i] + 1;$

If e is local:

$V_i(e) := V_i$

Else If e is send request m :

$V_i(e) := V_i$


Send (m, V_i)

Else If e is receive(m, t) event:

$V_i[j] := \max(V_i[j], t[j])$, for $j = 1, 2, \dots, N$.
(componentwise maximum)

COMP 512: Distributed Systems

Vector timestamps



COMP 512: Distributed Systems